

Bitcoin

50.037 Blockchain Technology
Paweł Szałachowski

Bitcoin

Bitcoin: A Peer-to-Peer Electronic Cash System

Satoshi Nakamoto
satoshin@gmx.com
www.bitcoin.org

Abstract. A purely peer-to-peer version of electronic cash would allow online payments to be sent directly from one party to another without going through a financial institution. Digital signatures provide part of the solution, but the main benefits are lost if a trusted third party is still required to prevent double-spending. We propose a solution to the double-spending problem using a peer-to-peer network. The network timestamps transactions by hashing them into an ongoing chain of hash-based proof-of-work, forming a record that cannot be changed without redoing the proof-of-work. The longest chain not only serves as proof of the sequence of events witnessed, but proof that it came from the largest pool of CPU power. As long as a majority of CPU power is controlled by nodes that are not cooperating to attack the network, they'll generate the longest chain and outpace attackers. The network itself requires minimal structure. Messages are broadcast on a best effort basis, and nodes can leave and rejoin the network at will, accepting the longest proof-of-work chain as proof of what happened while they were gone.

- White paper, 2008
- Open and distributed consensus (with incentives and no trusted parties)
- Peer-to-peer (equal participants)
- Cheap, anonymous/private, instant, and censorship-free payments
- Bitcoins (₿)
 - current supply: 17M₿
 - total supply: 21M₿
 - 100,000,000 Satoshi = 1₿

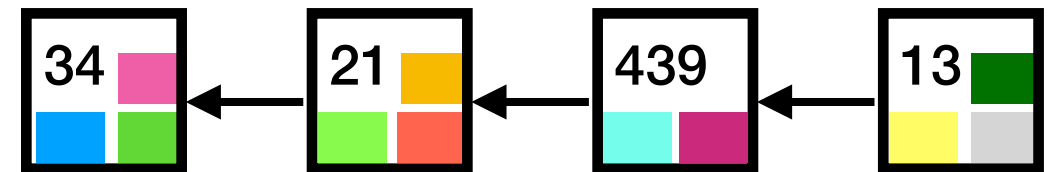
Bitcoin

- Previous systems
 - Ecash (D. Chaum)
 - Cypherpunks
- Issues
 - Centralization and Trust
 - Lack of incentives
 - Security
- **Bitcoin** (Core)
- Bitcoin Cash
- Bitcoin Gold
- Bitcoin Diamond
- Bitcoin Private
- ... hundreds forks ...

Recap & High-level Overview

Bitcoin's Blockchain

- Hash chain of blocks



- Block

- Set of transactions

- Forming hash tree

- Special number (Nonce)

- Hash pointer/link to the previous block

- Proof of work

- $H(\text{block}) < \text{target}$ (target defines difficulty)

- You can see leading zeros in hashes

- $H(.)$ defined as $\text{SHA256}(\text{SHA256}(.))$

$$H\left(\begin{array}{|c|c|} \hline 34 & \text{pink} \\ \hline \text{blue} & \text{green} \\ \hline \end{array}\right)$$

00000000892fcb17...

$$H\left(\begin{array}{|c|c|} \hline 21 & \text{yellow} \\ \hline \text{green} & \text{red} \\ \hline \end{array}\right)$$

00000000490747db...


$$H\left(\begin{array}{|c|c|} \hline 439 & \text{cyan} \\ \hline \text{cyan} & \text{magenta} \\ \hline \end{array}\right)$$

00000000a4d587e0...

$$H\left(\begin{array}{|c|c|} \hline 13 & \text{yellow} \\ \hline \text{yellow} & \text{grey} \\ \hline \end{array}\right)$$

00000000ff410ef45...

Transactions and Mining




Low-level Overview

Block and Header

- Block
 - Block header
 - Transactions (form a hash tree)


Field	Purpose
Version	Block version number
hashPrevBlock	256-bit hash of the previous block header
hashMerkleRoot	256-bit hash based on all of the transactions in the block
Time	Current timestamp as seconds since 1970-01-01T00:00 UTC
Bits	Current target in compact format
Nonce	32-bit number (starts at 0)




Proof of Work

- [illegible]

Bitcoin Hash Rate vs Difficulty (2 Months)



Bitcoin Block Generation Time vs Difficulty



Block and Header Validation

- Miners propagate blocks
- A new block is accepted if
 - Size \leq MAX_BLOCK_SIZE (1MB, very controversial topic...)
 - Hash of its header $<$ current target
 - The previous block's hash is correct
 - Transactions form a hash tree whose root is in the header
 - All its transactions are *correct*
 - The timestamps does not deviate from the network time
 - ...

Full Nodes

- Full Nodes
 - Replicated state machine
 - **Share state** managed by replicas
 - Set of unspent coins
 - Apply **operations** to modify it over time
 - Transactions
 - Receive, validate, and store all blocks
- Miners are full nodes running the consensus protocol
 - Propagate new blocks

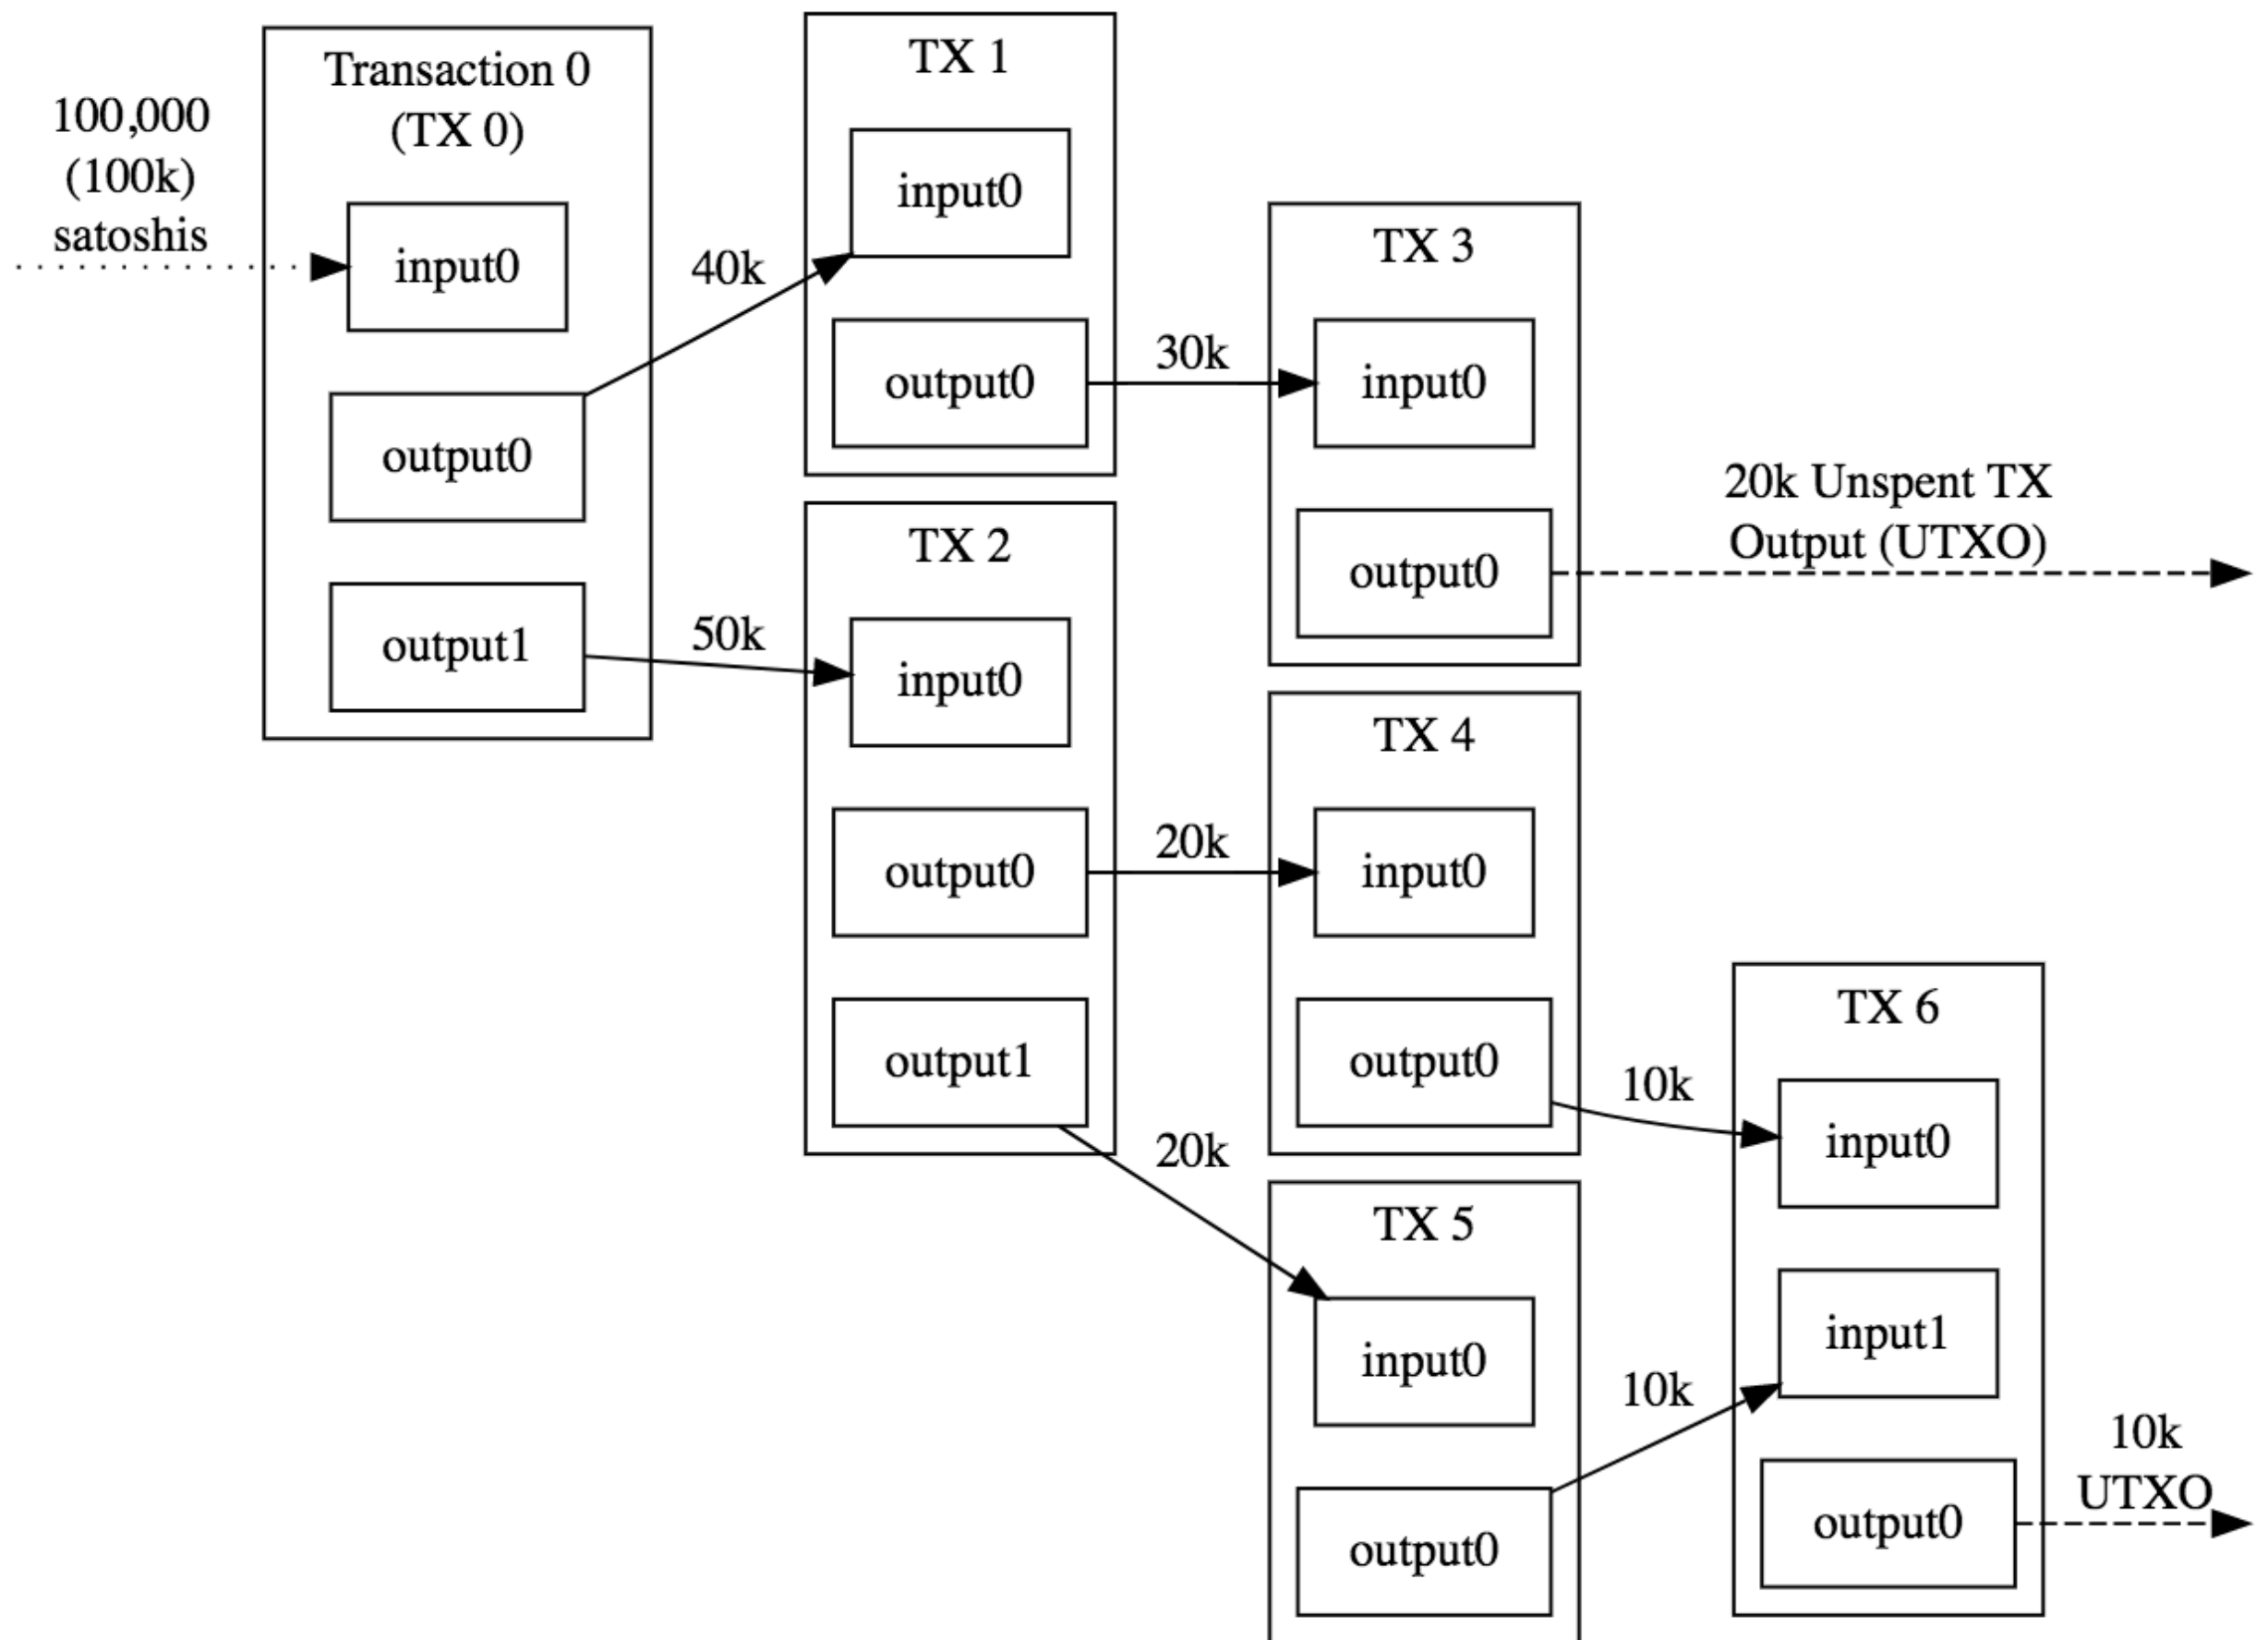
Transactions and the UTXO Model

Transaction Data

- Transaction
 - Sender (coins to be spent)
 - Receiver (how to spend them later)
 - Fee (incentivize miners)
- One or more transaction per block
- Transactions are identified by their hashes
 - $\text{TXID} = H(\text{transaction})$

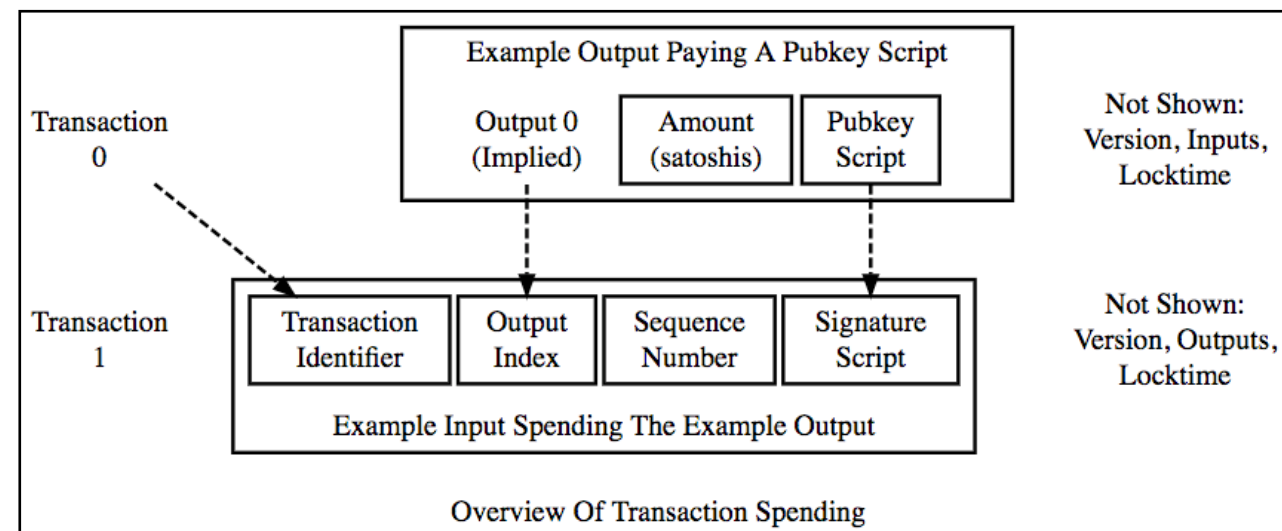
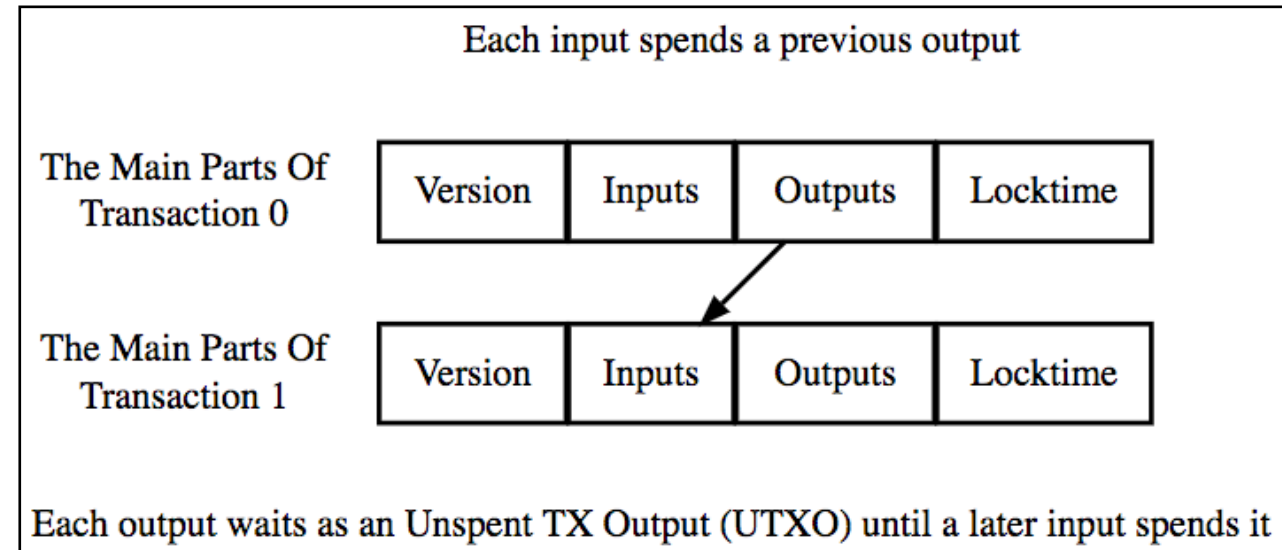
UTXO Transaction Model

- The Unspent Transaction Output (UTXO) Model
- Each transaction spends the coins previously received in one or more earlier transactions
 - The input of one transaction is the output of a previous transaction
- A single transaction can create multiple outputs
 - Each output can only be used as an input once
- Transaction validity
 - Must only use UTXOs as inputs
 - Value of its outputs \leq inputs
 - The difference (inputs - outputs) can be claimed as a fee by the miner who creates the block containing this transaction



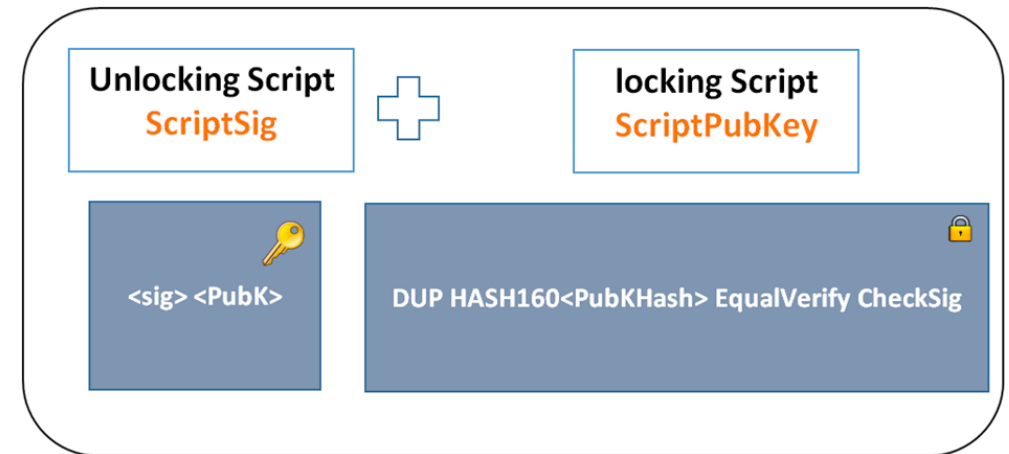
Transaction Format

- Version determines set of rules to apply
- An input uses a TXID and an output index number to identify a particular output to be spent
- It also has a **signature script** which allows it to provide data parameters that satisfy the conditions (**pubkey script**) to spend the output
- Locktime
 - The earliest time a transaction can be added to the blockchain
 - New non-locked transaction can invalidate it (use the same outputs




Pay-To-Public-Key-Hash (P2PKH)

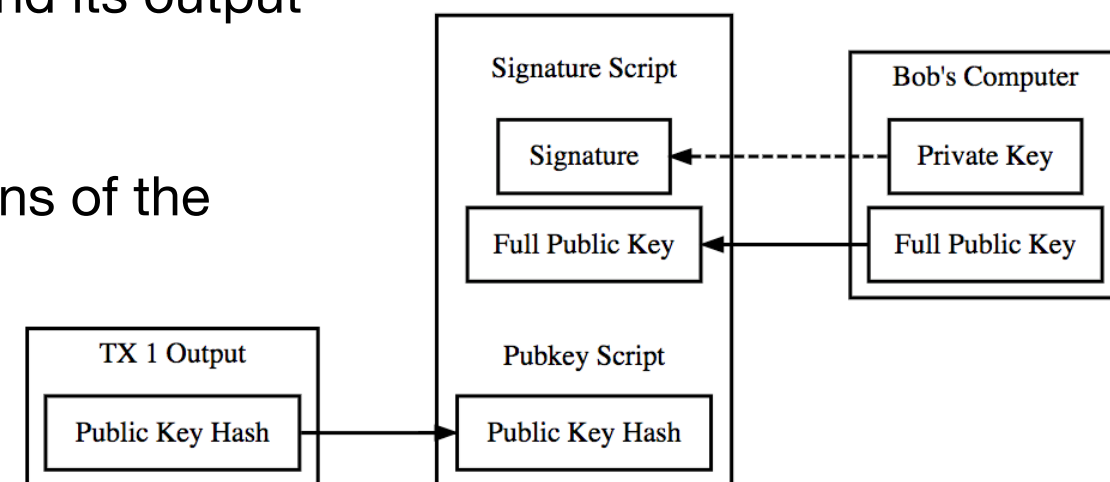
- Digital signatures (public/private keys)
- Alice sends transaction to Bob's address
 - Address is encoded as public key's hash (base58)
- P2PKH transaction output that allows to spend it anyone who can prove its ownership of the private key corresponding to Bob's hashed public key
 - These instructions are called the **pubkey script** or scriptPubKey
- Transaction is propagated and Bob treats it as a spendable balance
- Bob spends it creating an input referring to the Alice's TXID and its output index
 - and creating a **signature script** that satisfies the conditions of the output's pubkey script



src: cryptocompare.com



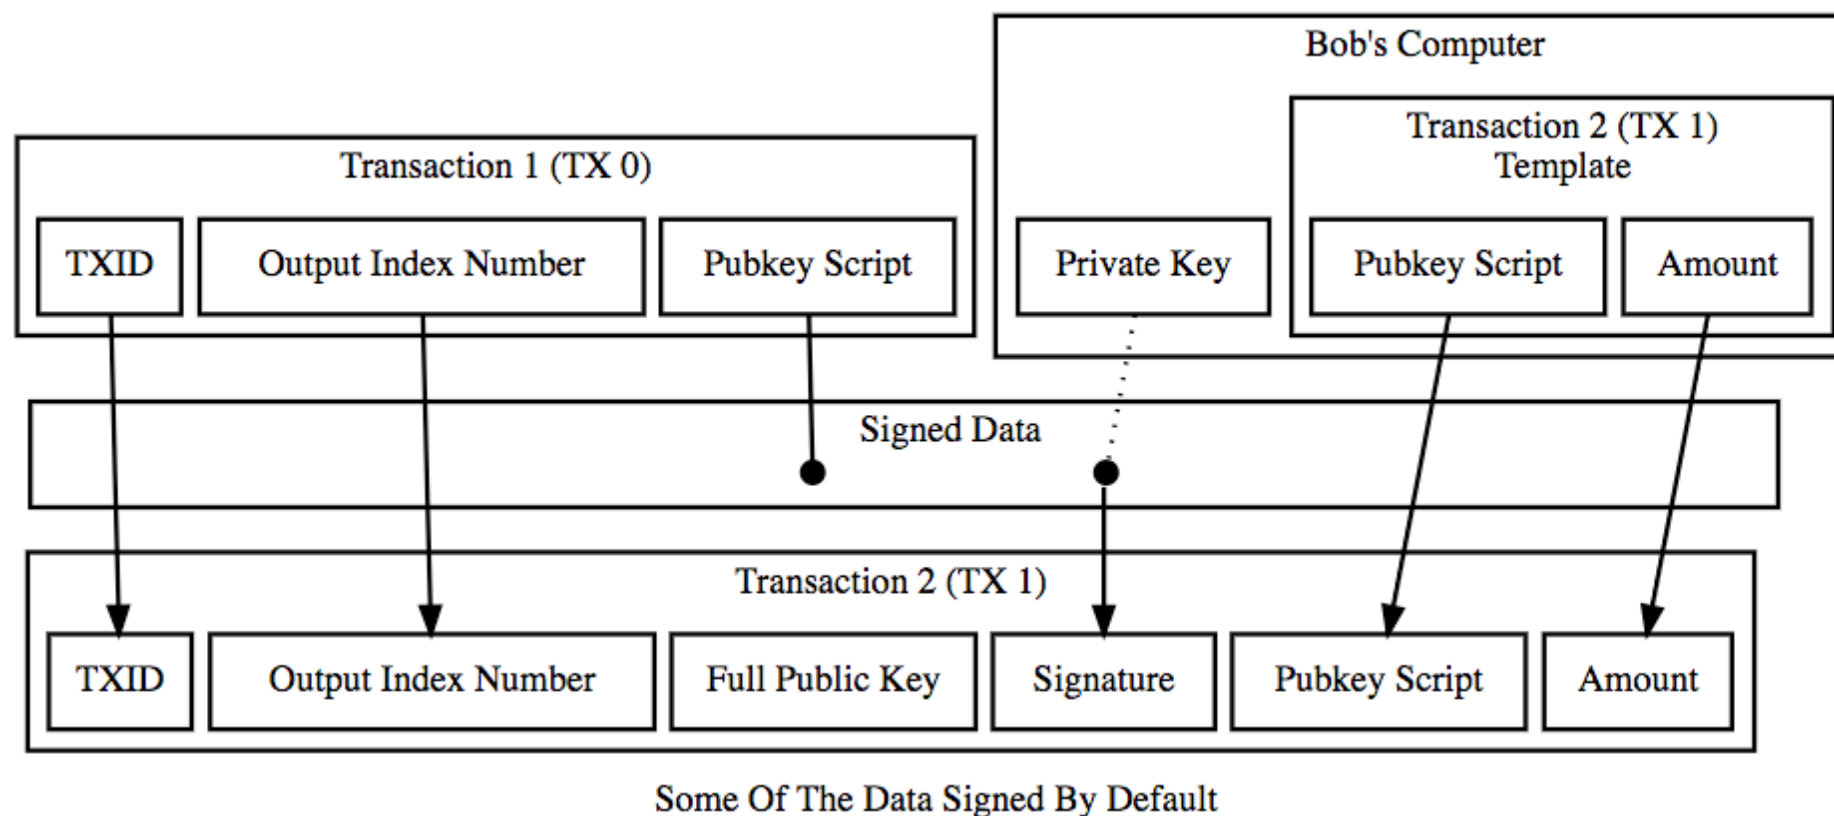
Creating A P2PKH Public Key Hash To Receive Payment



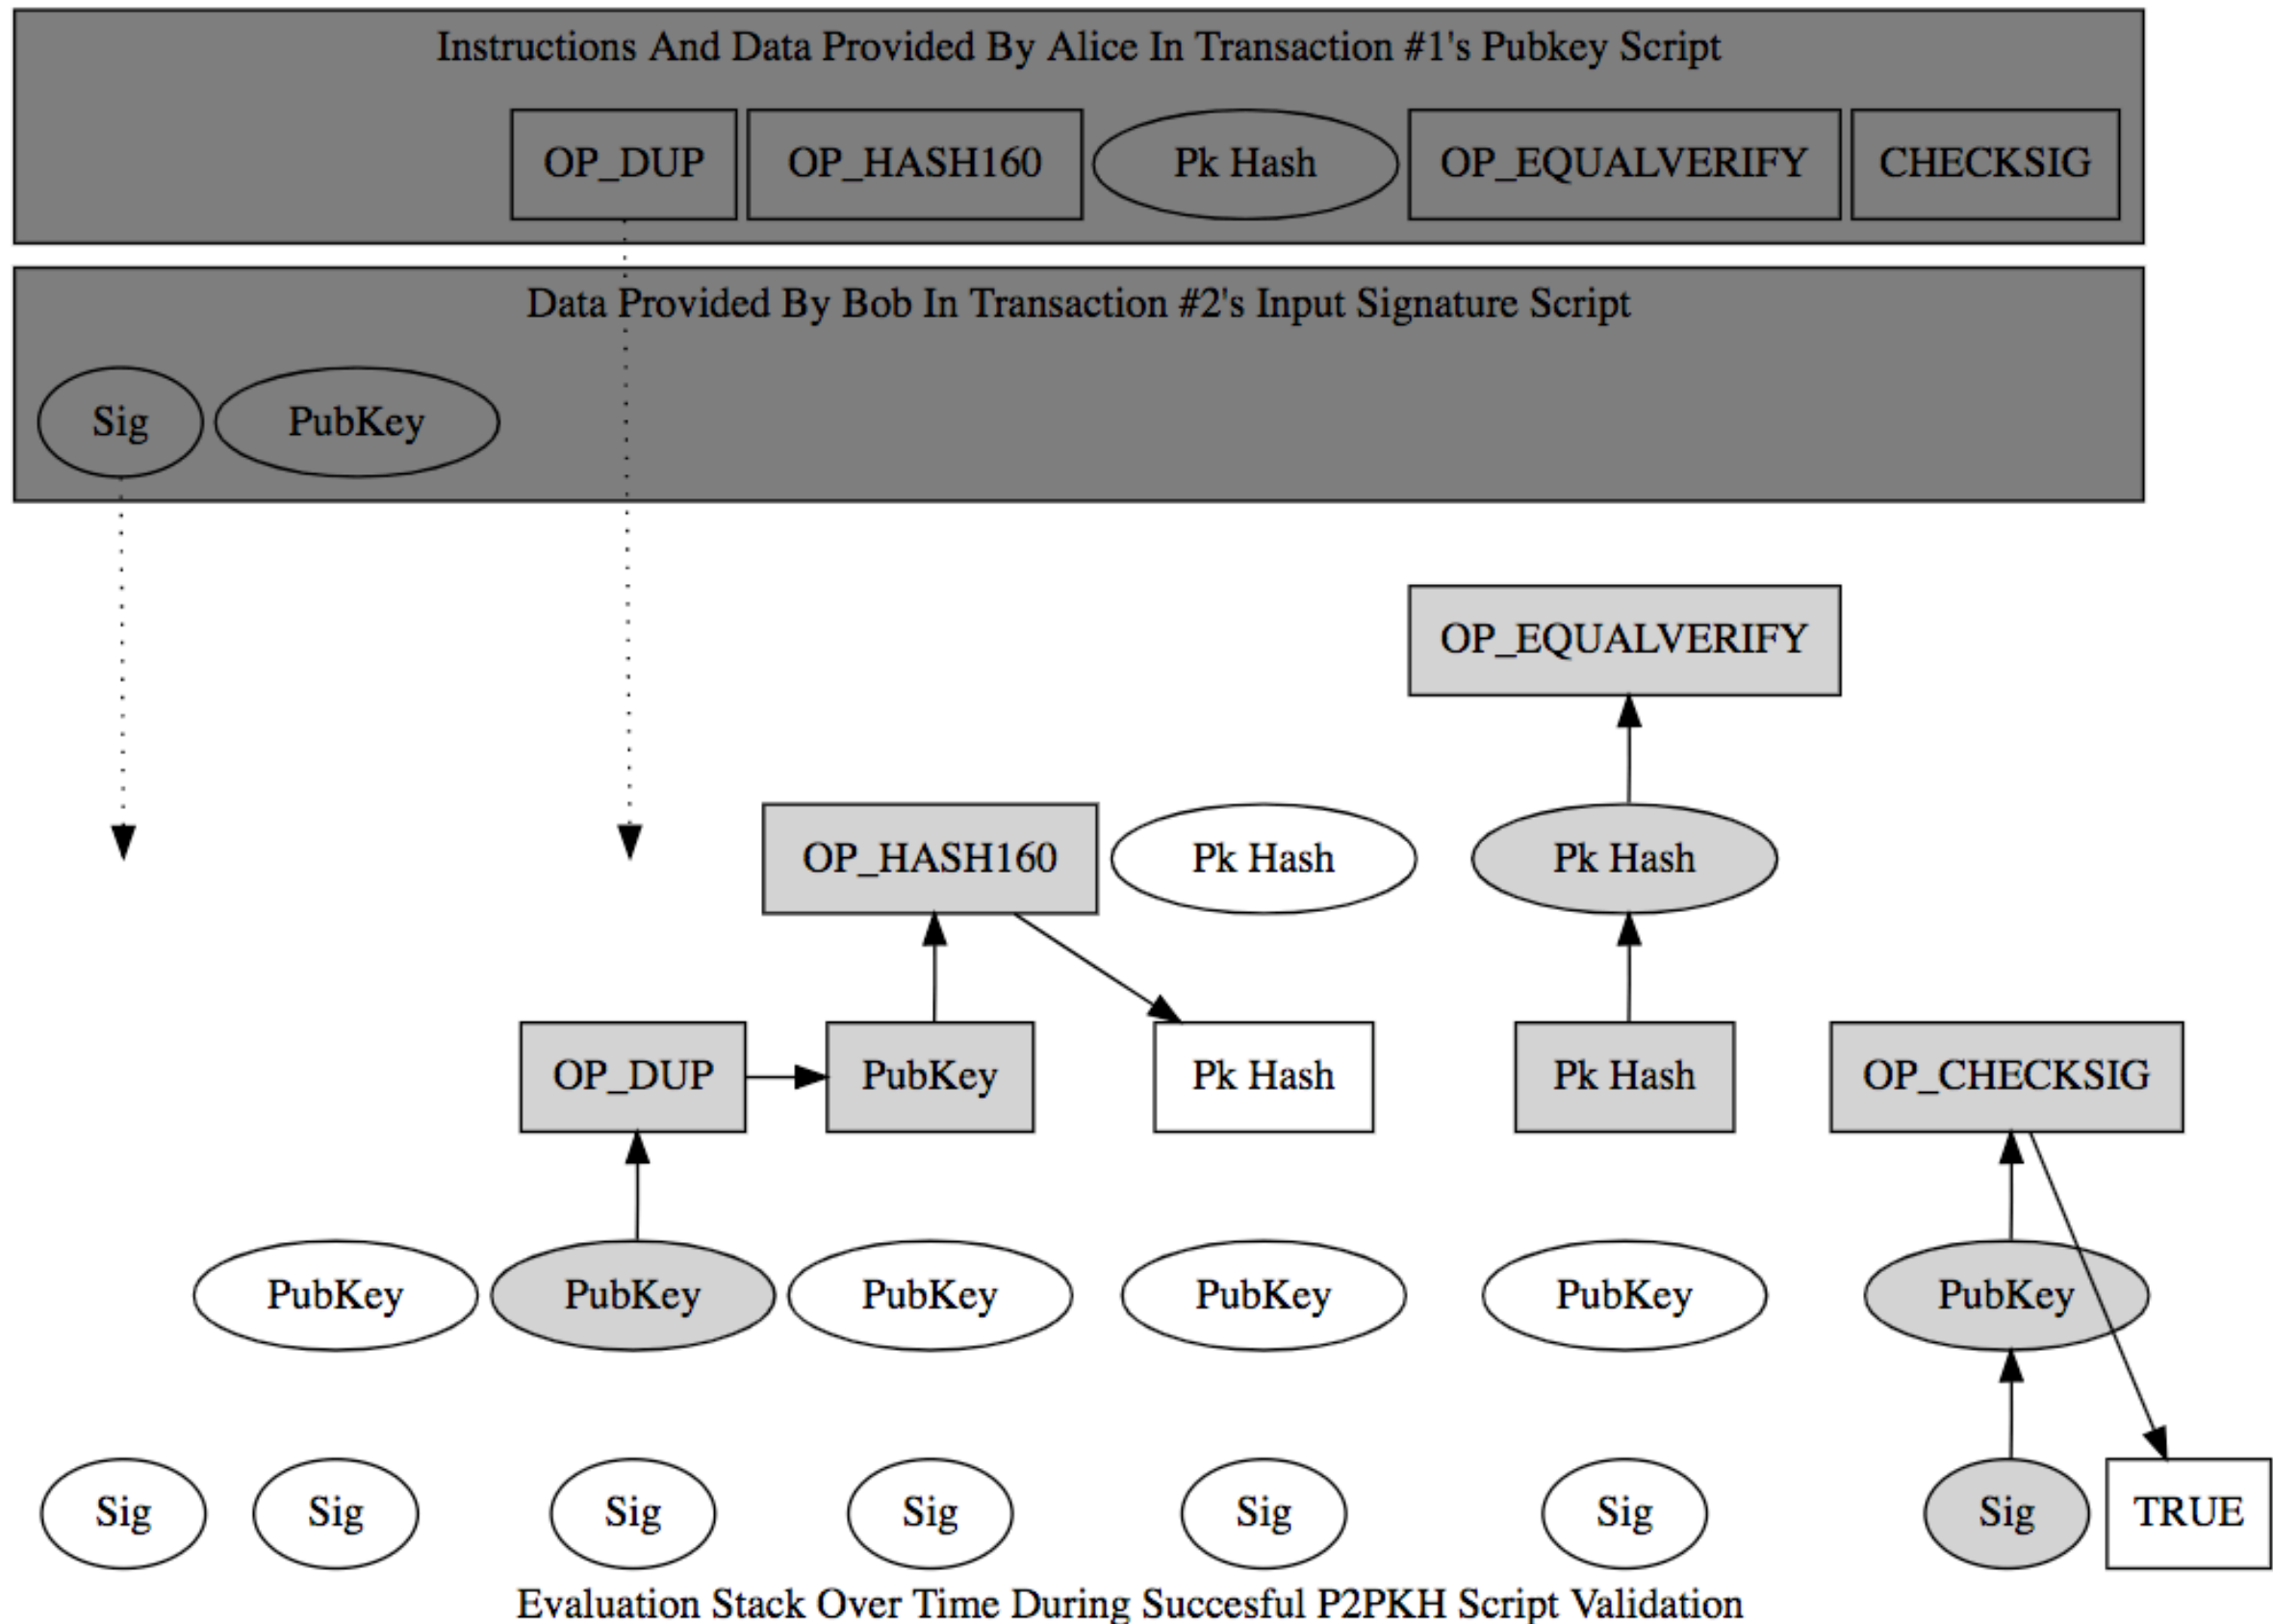
Spending A P2PKH Output

P2PKH

- Bob's signature script
 - His public key (script can check its hash)
 - Signature (proves possession of the corresponding private key and makes the transaction secure)



P2PKH Script Validation




P2PKH Example



Other Scripts

- Pay-To-Script-Hash (P2SH)



Creating A P2SH Redeem Script Hash To Receive Payment

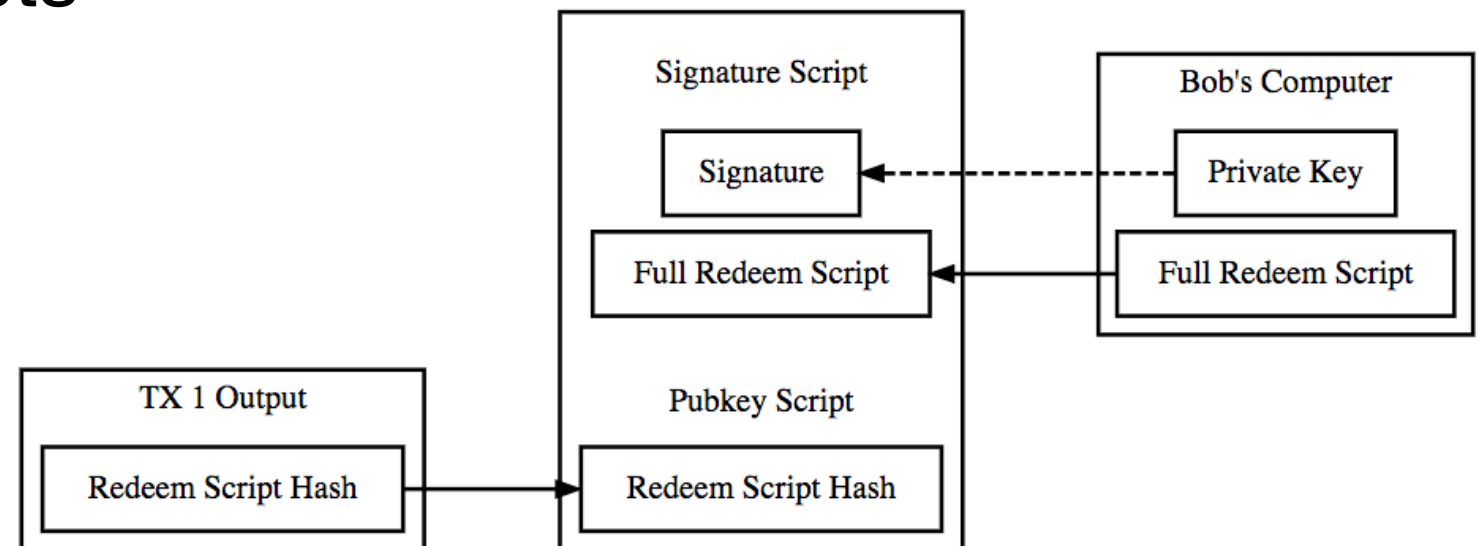
- (Almost) arbitrary spending rules

- Possible to create contracts

- Multisig (m-of-n)

- Escrow

- ...



Spending A P2SH Output

- Only some of them are supported in practice

Transactions (Misc)

- OP_RETURN <data>
 - Returns <data>, cannot be spent
- Fees
 - Price per byte
 - Prioritize transactions
 - Change outputs
 - Mistaken transaction
 - No change output specified

Transaction View information about a bitcoin transaction

258478e8b7a3b78301661e78b4f93a792af878b545442498065ab272eaacf035

1LtjWsKsrr2RweDLAmv75oGL7tjVF4wx7W → 1CfsAiYaVfk12dnZpZALcRSP9jjWDk26FX 0.01252199 BTC
1CfsAiYaVfk12dnZpZALcRSP9jjWDk26FX

0.01252199 BTC

Summary

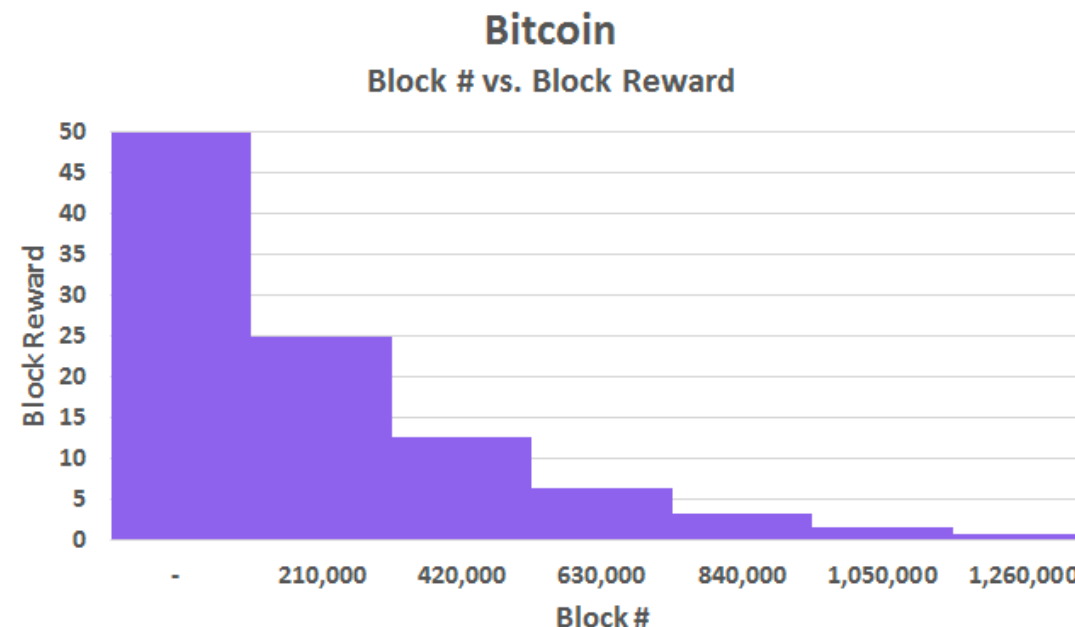
Size	341 (bytes)
Weight	1364
Received Time	2013-09-17 21:20:13
Included In Blocks	258546 (2013-09-17 21:23:26 + 3 minutes)
Confirmations	284269
Visualize	View Tree Chart

Inputs and Outputs

Total Input	80.99252199 BTC
Total Output	0.01252199 BTC
Fees	80.98 BTC
Fee per byte	23,747,800.587 sat/B
Fee per weight unit	5,936,950.147 sat/WU
Estimated BTC Transacted	0 BTC
Scripts	Show scripts & coinbase

Coin Minting

- But how coins are actually introduced to the system?
- **Coinbase transactions** (the first transaction of every block)
 - Collecting and spending block reward plus transaction fees paid by transactions of this block
 - Coinbase coins cannot be spent for at least 100 blocks
 - Prevent immediate spending blocks that may turn out to be stale
- Reward scheme



Genesis Block


- Input field of a coinbase transaction can be an arbitrary text

00000000	01 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00
00000010	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00
00000020	00 00 00 00 3B A3 ED FD	7A 7B 12 B2 7A C7 2C 3E;fíýz{.²zÇ,>
00000030	67 76 8F 61 7F C8 1B C3	88 8A 51 32 3A 9F B8 AA	gv.a.È.Ã~ŠQ2:Ÿ_a
00000040	4B 1E 5E 4A 29 AB 5F 49	FF FF 00 1D 1D AC 2B 7C	K.^J)«_Iÿÿ...¬+
00000050	01 01 00 00 00 01 00 00	00 00 00 00 00 00 00 00
00000060	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00
00000070	00 00 00 00 00 00 FF FF	FF FF 4D 04 FF FF 00 1DÿÿÿÿM.ÿÿ..
00000080	01 04 45 54 68 65 20 54	69 6D 65 73 20 30 33 2F	..EThe Times 03/
00000090	4A 61 6E 2F 32 30 30 39	20 43 68 61 6E 63 65 6C	Jan/2009 Chancel
000000A0	6C 6F 72 20 6F 6E 20 62	72 69 6E 6B 20 6F 66 20	lor on brink of
000000B0	73 65 63 6F 6E 64 20 62	61 69 6C 6F 75 74 20 66	second bailout f
000000C0	6F 72 20 62 61 6E 6B 73	FF FF FF FF 01 00 F2 05	or banksÿÿÿÿ..ð.
000000D0	2A 01 00 00 00 43 41 04	67 8A FD B0 FE 55 48 27	*....CA.gŠý°pUH'
000000E0	19 67 F1 A6 71 30 B7 10	5C D6 A8 28 E0 39 09 A6	.gñ q0•.\Ö" (à9.
000000F0	79 62 E0 EA 1F 61 DE B6	49 F6 BC 3F 4C EF 38 C4	ybaê.aþ¶IÖ¼?Li8Ä
00000100	F3 55 04 E5 1E C1 12 DE	5C 38 4D F7 BA 0B 8D 57	óU.â.Á.þ\8M÷ø..W
00000110	8A 4C 70 2B 6B F1 1D 5F	AC 00 00 00 00	ŠLp+kñ._¬....



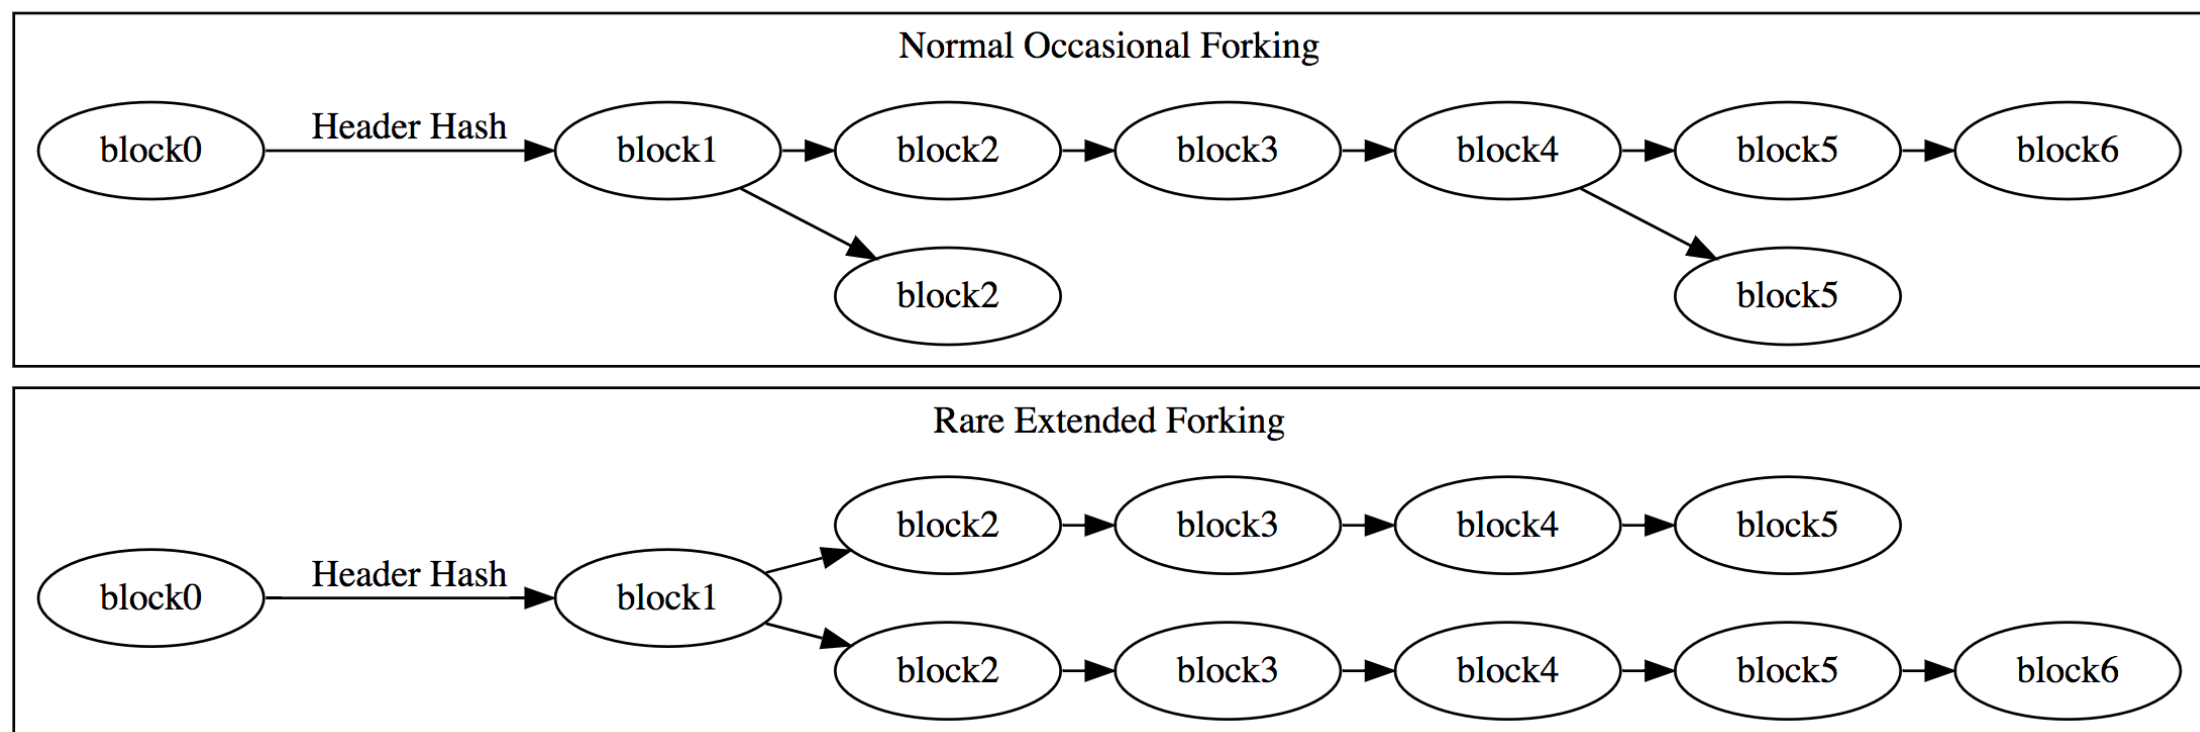
Forks

Forks (Recap)

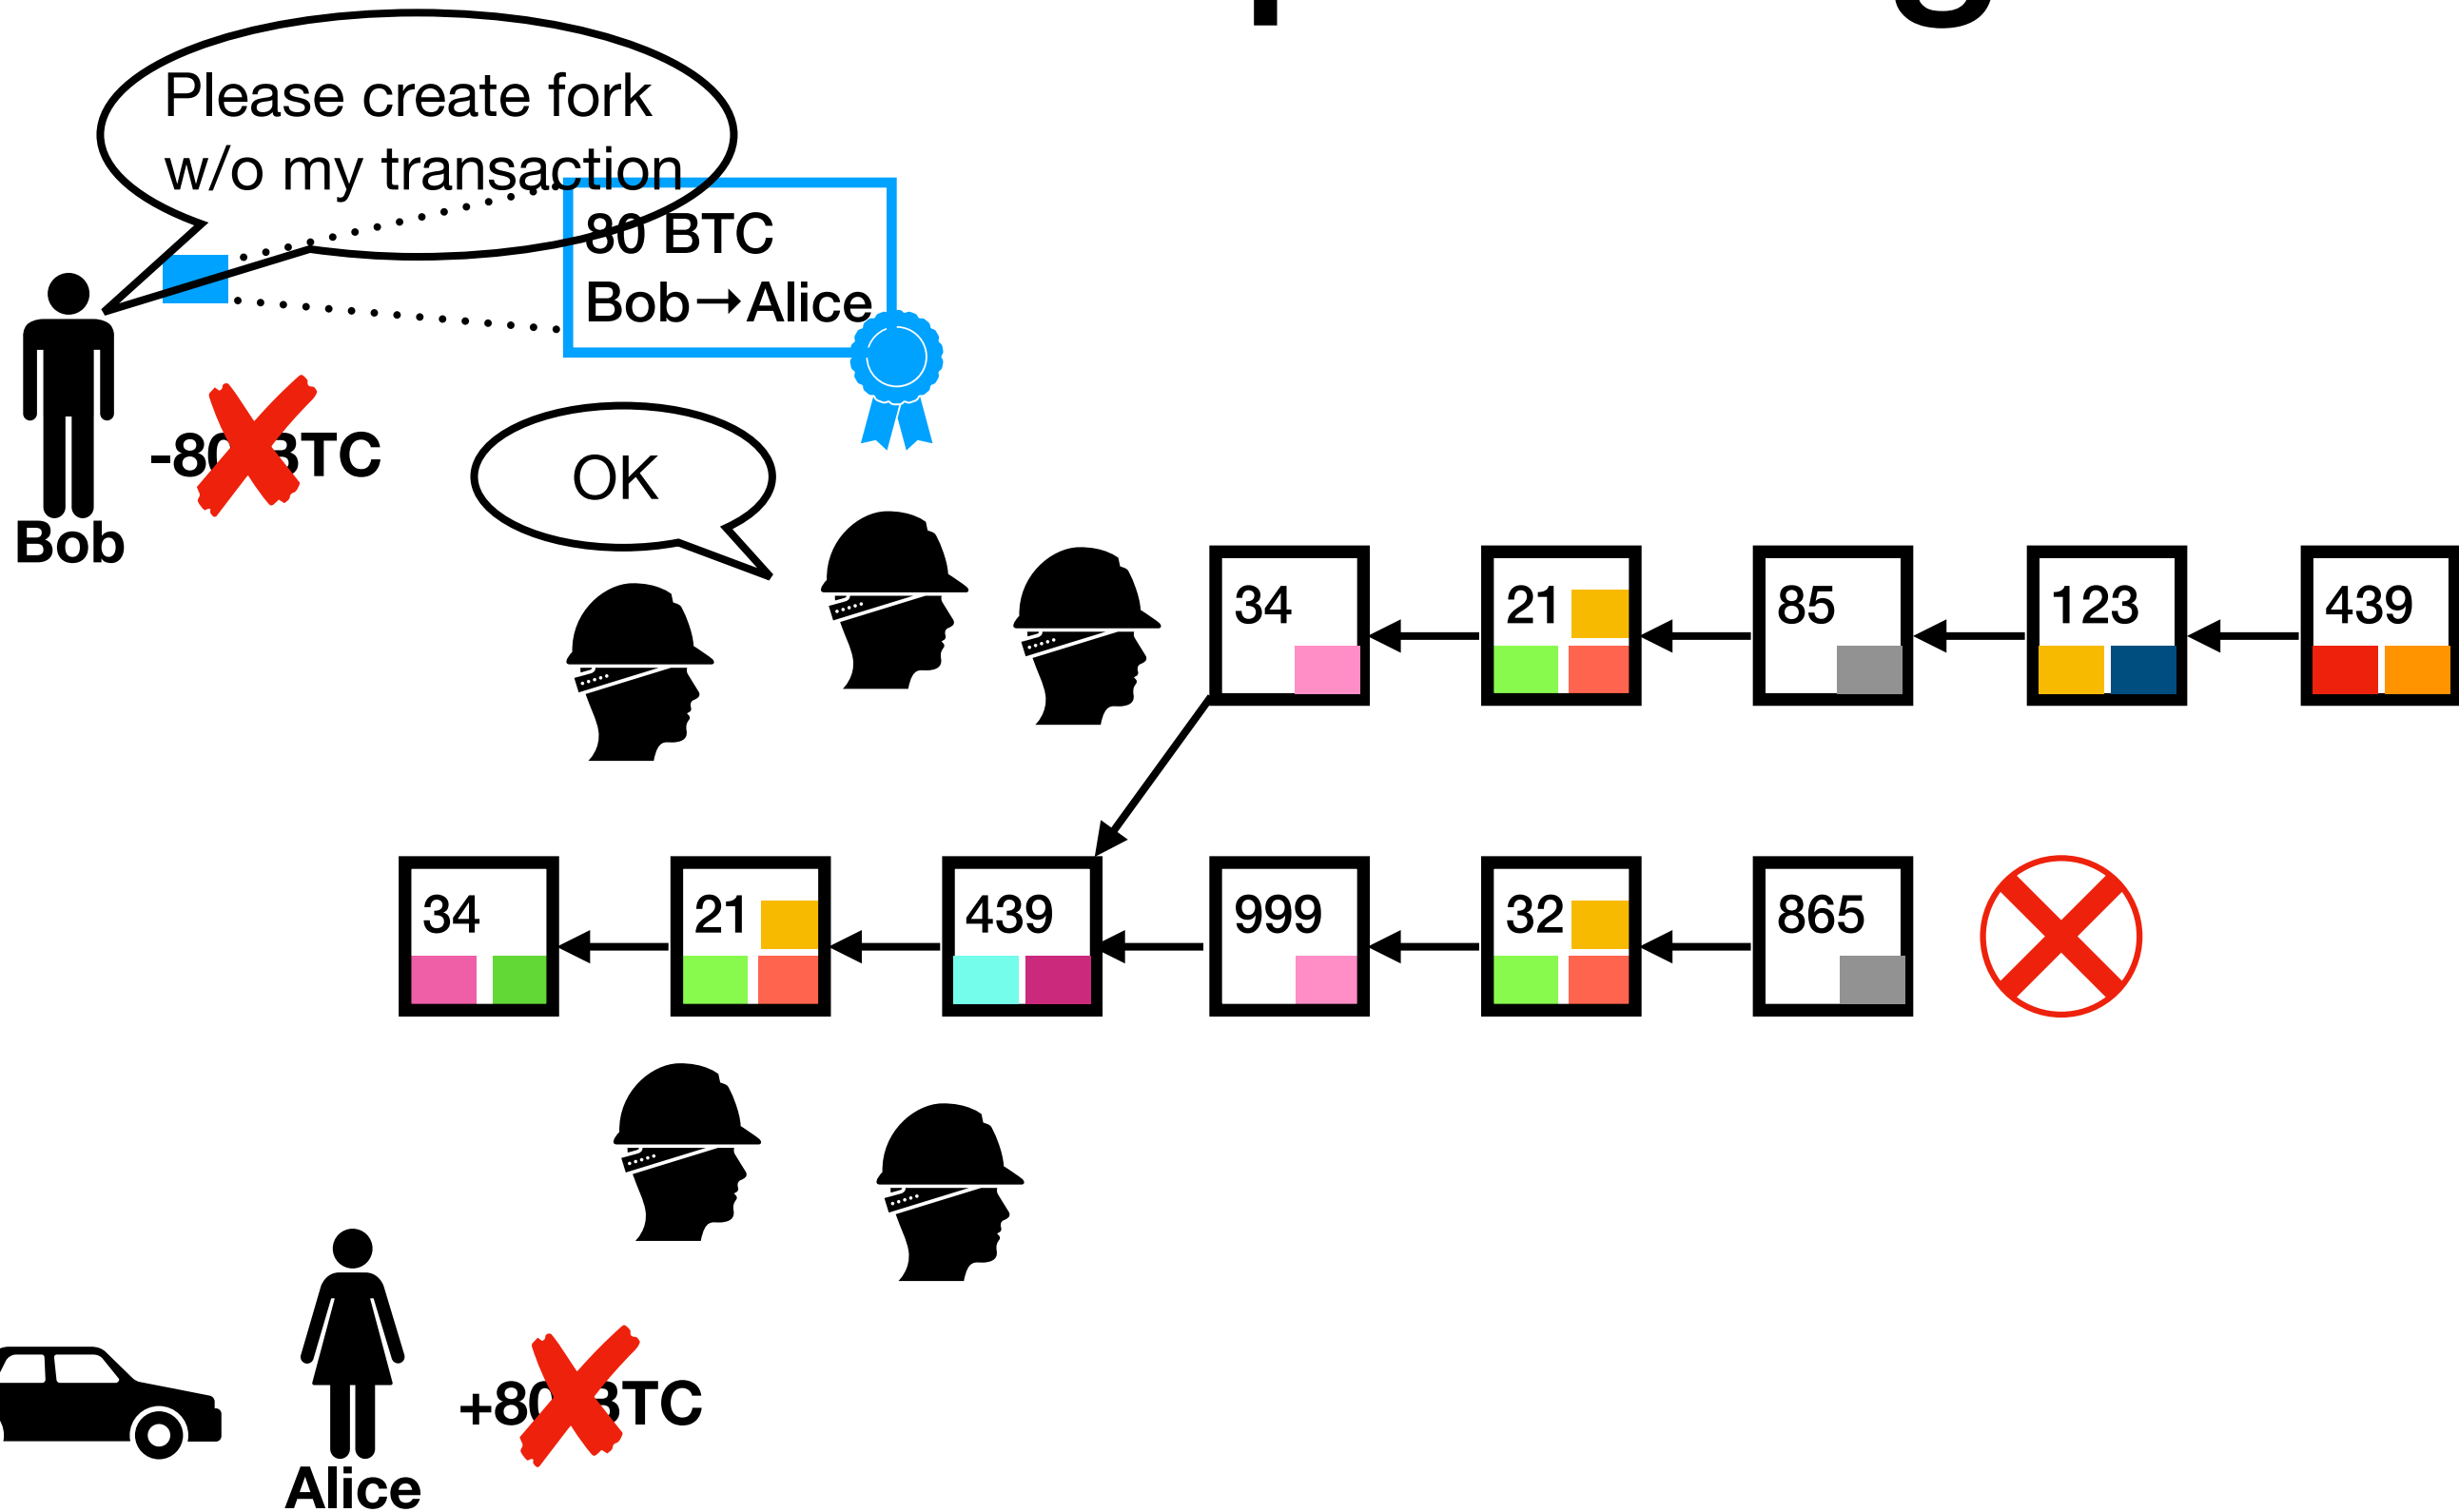


Forks

- In the case of one-block fork, nodes individually decide which to accept (usually, the first seen)
- Nodes follow the strongest chain and discard stale blocks of weaker chains
 - Sometimes these stale blocks are (incorrectly) called orphans
- Why forks are bad?



Double Spending


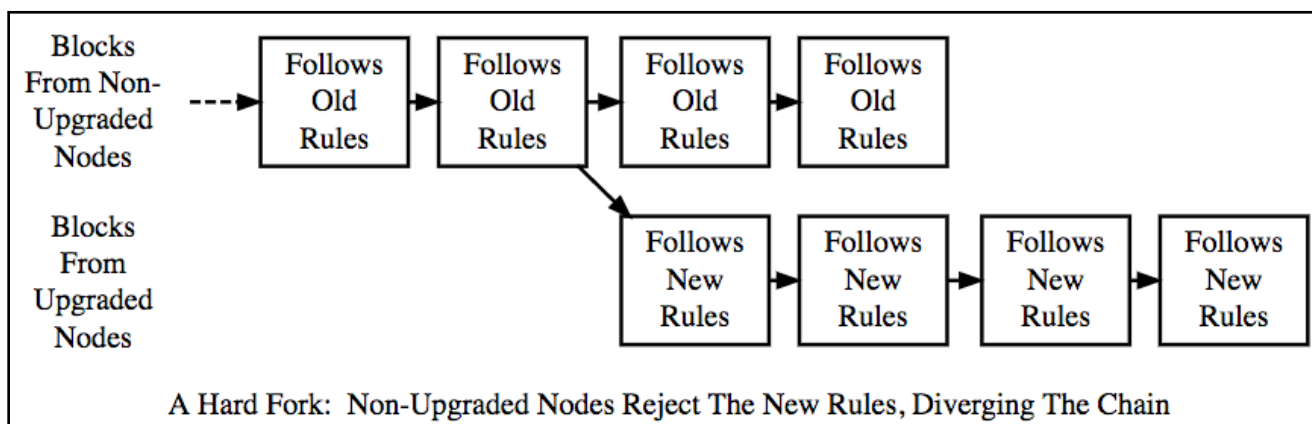


Double Spending Protection

- Confirmations
 - With honest majority of hash power long forks are unlikely
 - 6 confirmation blocks suggested (~1h)
- Software protection
 - Checkpoints implemented
 - (actually, that is a protocol violation)

Consensus Rule Changes

- Nodes use the same consensus rules (replicated state machine), so what about updates (new features, bug fixes)?
- Updated and non-updated nodes will follow different rules
- Hard fork: block following the new consensus rules is accepted by upgraded but rejected by non-upgraded nodes
 - Chains will be divergent, nodes will never agree
- Soft fork: block violating the new consensus rules is rejected by upgraded but accepted by non-upgraded nodes
 - Chains will converge, nodes will agree (when majority of the hash rate applies new rules)
- Examples: changing block size, rejecting malformed transactions, ...
- Waiting for hash power, flag days, user/miner activated soft forks



Reading

- Textbook 2,3,4
- <https://bitcoin.org/en/developer-guide>
- <https://queue.acm.org/detail.cfm?id=3136559>