Scalability

50.037 Blockchain Technology

Paweł Szałachowski

Scalability

- “*Scalability is the capability of a system, network, or process to handle a growing amount of work, or its potential to be enlarged to accommodate that growth.*”
- Horizontal and vertical scaling

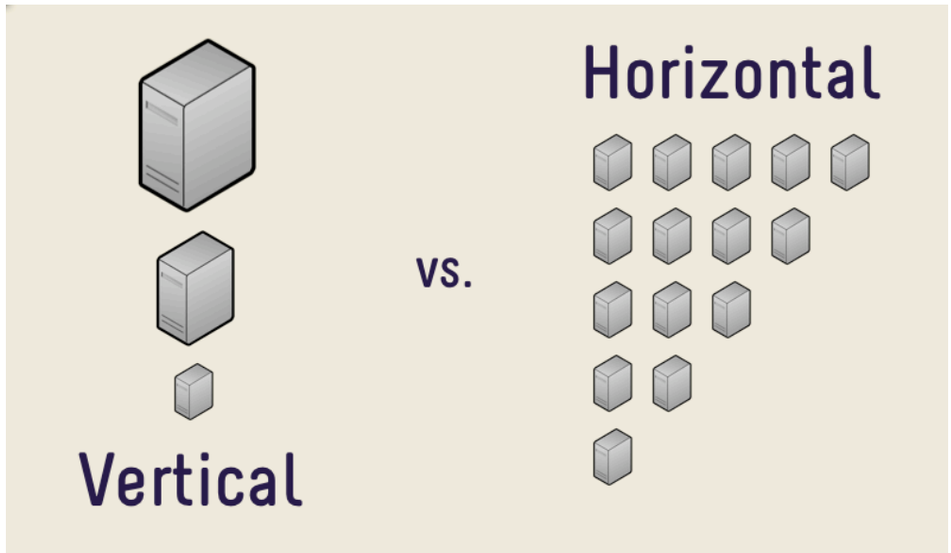

https://pudgylogic.blogspot.com/

Blockchain Stack

- Network: propagate transactions
  - Latency, bandwidth, # of nodes, …
- **Consensus**: order Txs
  - # of nodes, # of Txs (throughput)
- RSM: Txs validation, contract execution, …
  - State size, execution complexity, …
- Apps: use the current state to implement some logic

Current State

- Throughput
  - Bitcoin: 7 tx/s
  - Ethereum: 10 tx/s
  - Visa: 50k tx/s
- Strategy
  - Faster tx processing
  - Faster consensus
  - Parallel execution

Why needed?

- Adoption
  - Real-world apps require high throughput and low latency
- Inverse scale effect
  - Fees

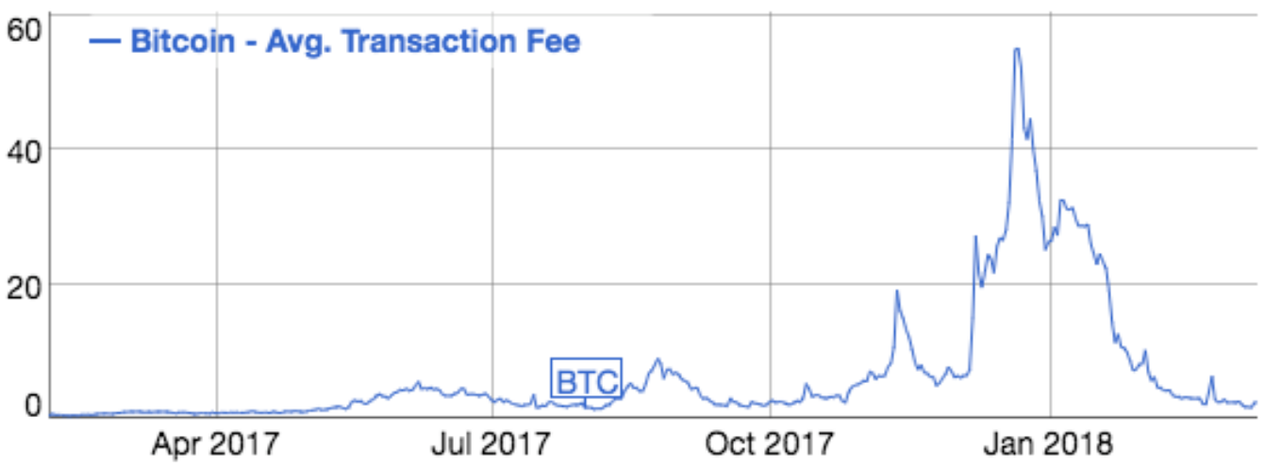

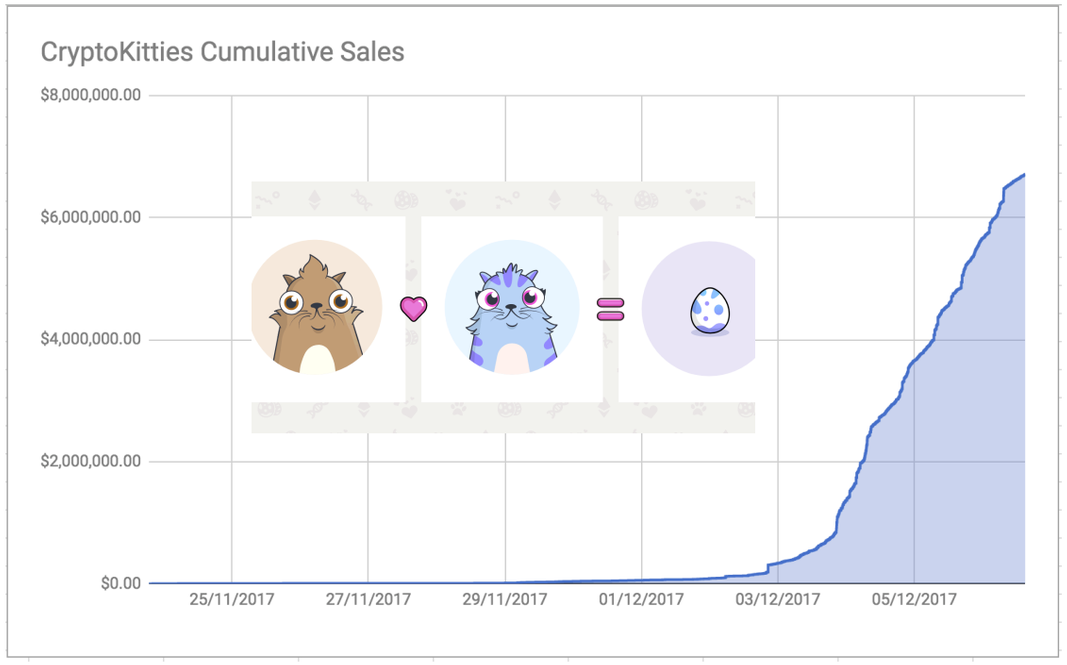

Blockchain Performance

- Bandwidth:
    - How many Txs can be processed?
- Latency
    - What is the consensus delay?
- Mining power utilization
    - The ratio between the mining power of the current chain and the mining power of the entire blockchain (describes stale block rate too), describes *security*
- Fairness
    - A miner should benefit from rewards proportionally to its mining power

Naive Improvements

- Blocks not every 10 minutes, but e.g., every 10 seconds
    - More forks => less mining power utilization => weaker security
- Larger blocks (very controversial topic BTW)
    - It takes longer to propagate
        - More forks => ….
    - Bitcoin -> Bitcoin Cash -> Bitcoin ABC vs Bitcoin SV

Security vs Performance

- Seems like Nakomoto consensus has some inherent tradeoffs
    - Security vs performance tradeoff
- Does it have to be like that?
    - We cannot significantly increase block or make them very frequent
- Design space
    - Why we need PoW?
    - Does it have to be combined with transactions propagation?

Bitcoin-NG

- https://www.usenix.org/node/194907 (paper, slides, talk)
- Insights
  - In Bitcoin, leader election and transaction serialization is combined
  - Why do not try to decouple it?
    - Elect leader via PoW and let her commit transactions
    - (Different order than in Bitcoin)

Bitcoin-NG

- Key blocks
  - Used for PoW-based leader election, i.e., H(header) < T
  - Point to the previous block (key or microblock)
    - The strongest-chain rule
- Microblocks
  - Generated by leader, at a defined rate
  - Contain header (with PrevHash) and a set of transactions

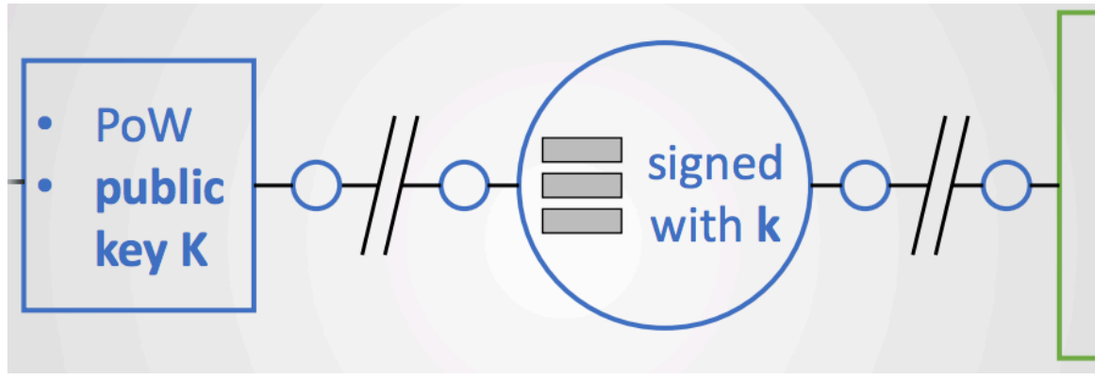

Bitcoin-NG

- Incentives
  - Leaders get rewards and tx fees
  - The next leader gets 60% of the previous tx fees (why?)

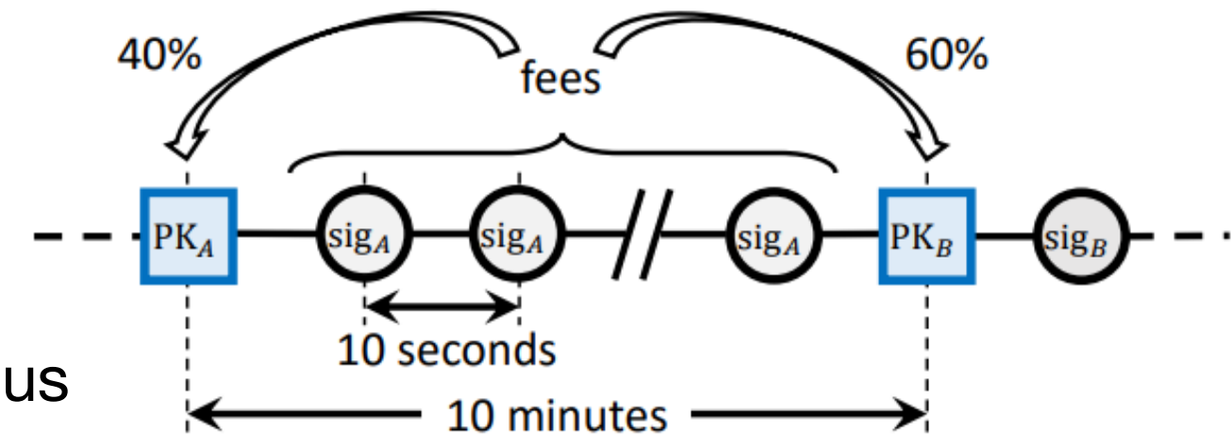

- Confirmations
  - Short forks will be frequent
- Microblock forks may be malicious
  - Entry with a proof of fraud can invalidate the revenue of malicious leaders

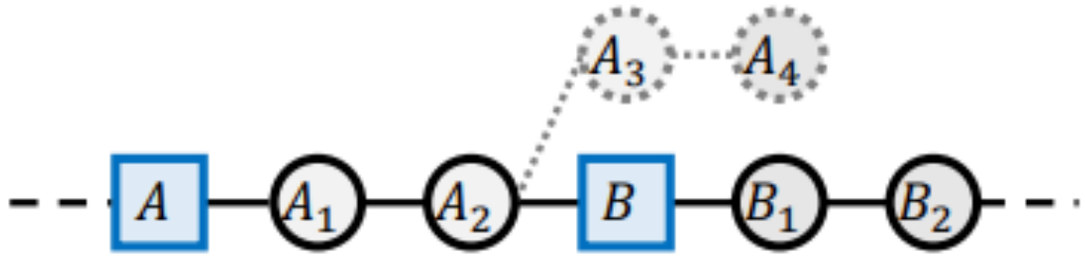

Bitcoin-NG

- Much better scalability and performance than Bitcoin
- Many systems build on this or similar ideas
- Everyone validates Txs
  - Throughput limited by a single machine
  - Can we do better?

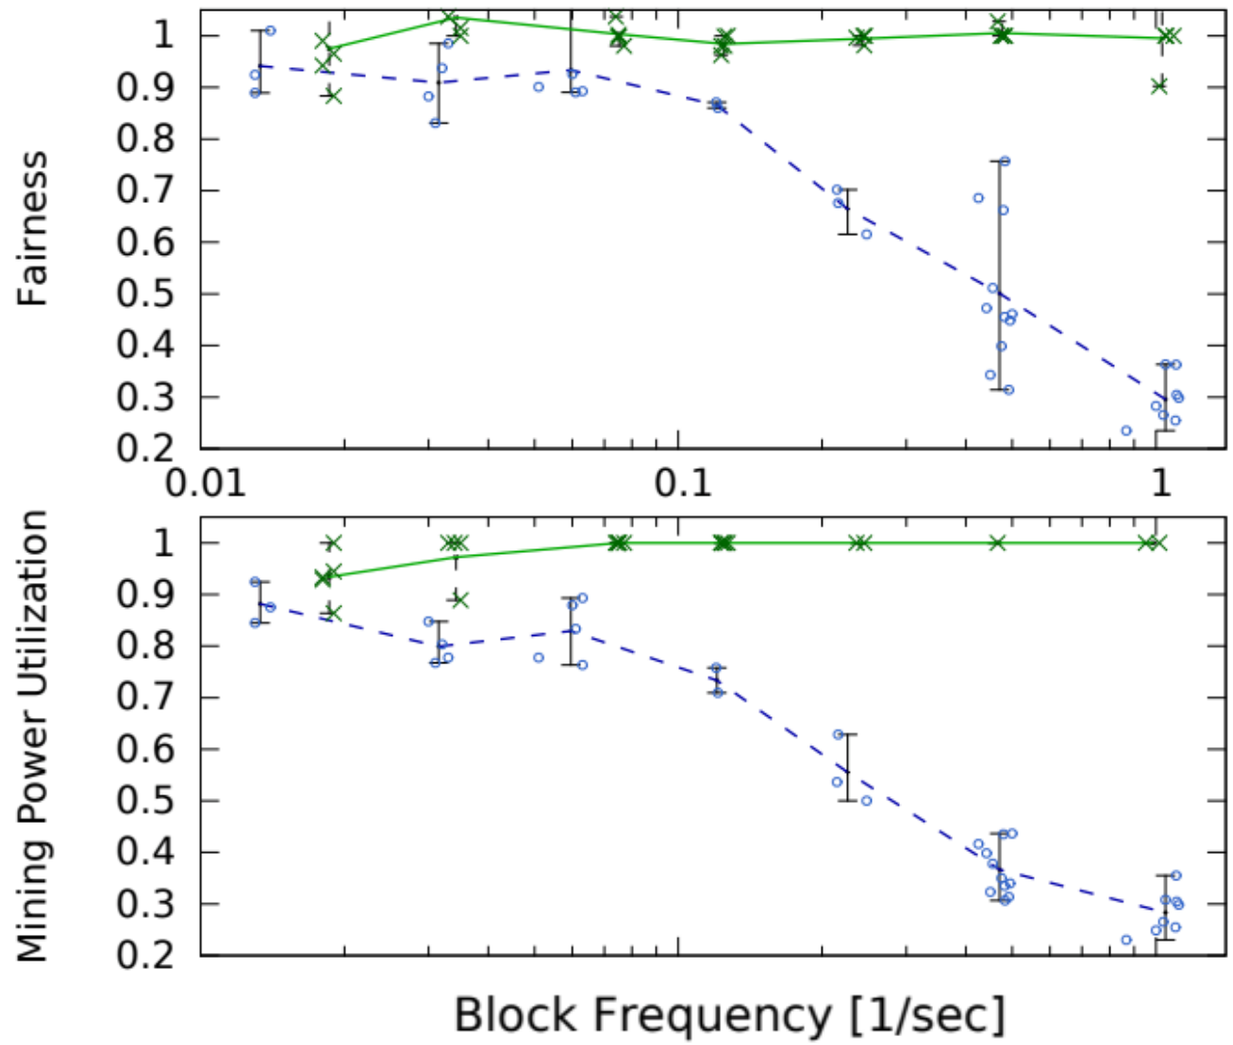

Sharding

Sharding

- The concept from database processing
- Divide transactions into groups and let different nodes process them
- Horizontal scaling
  - Throughput increases linearly as the network grows
- Ideas: establish identities via PoW, divide work, run BFT

Sharding

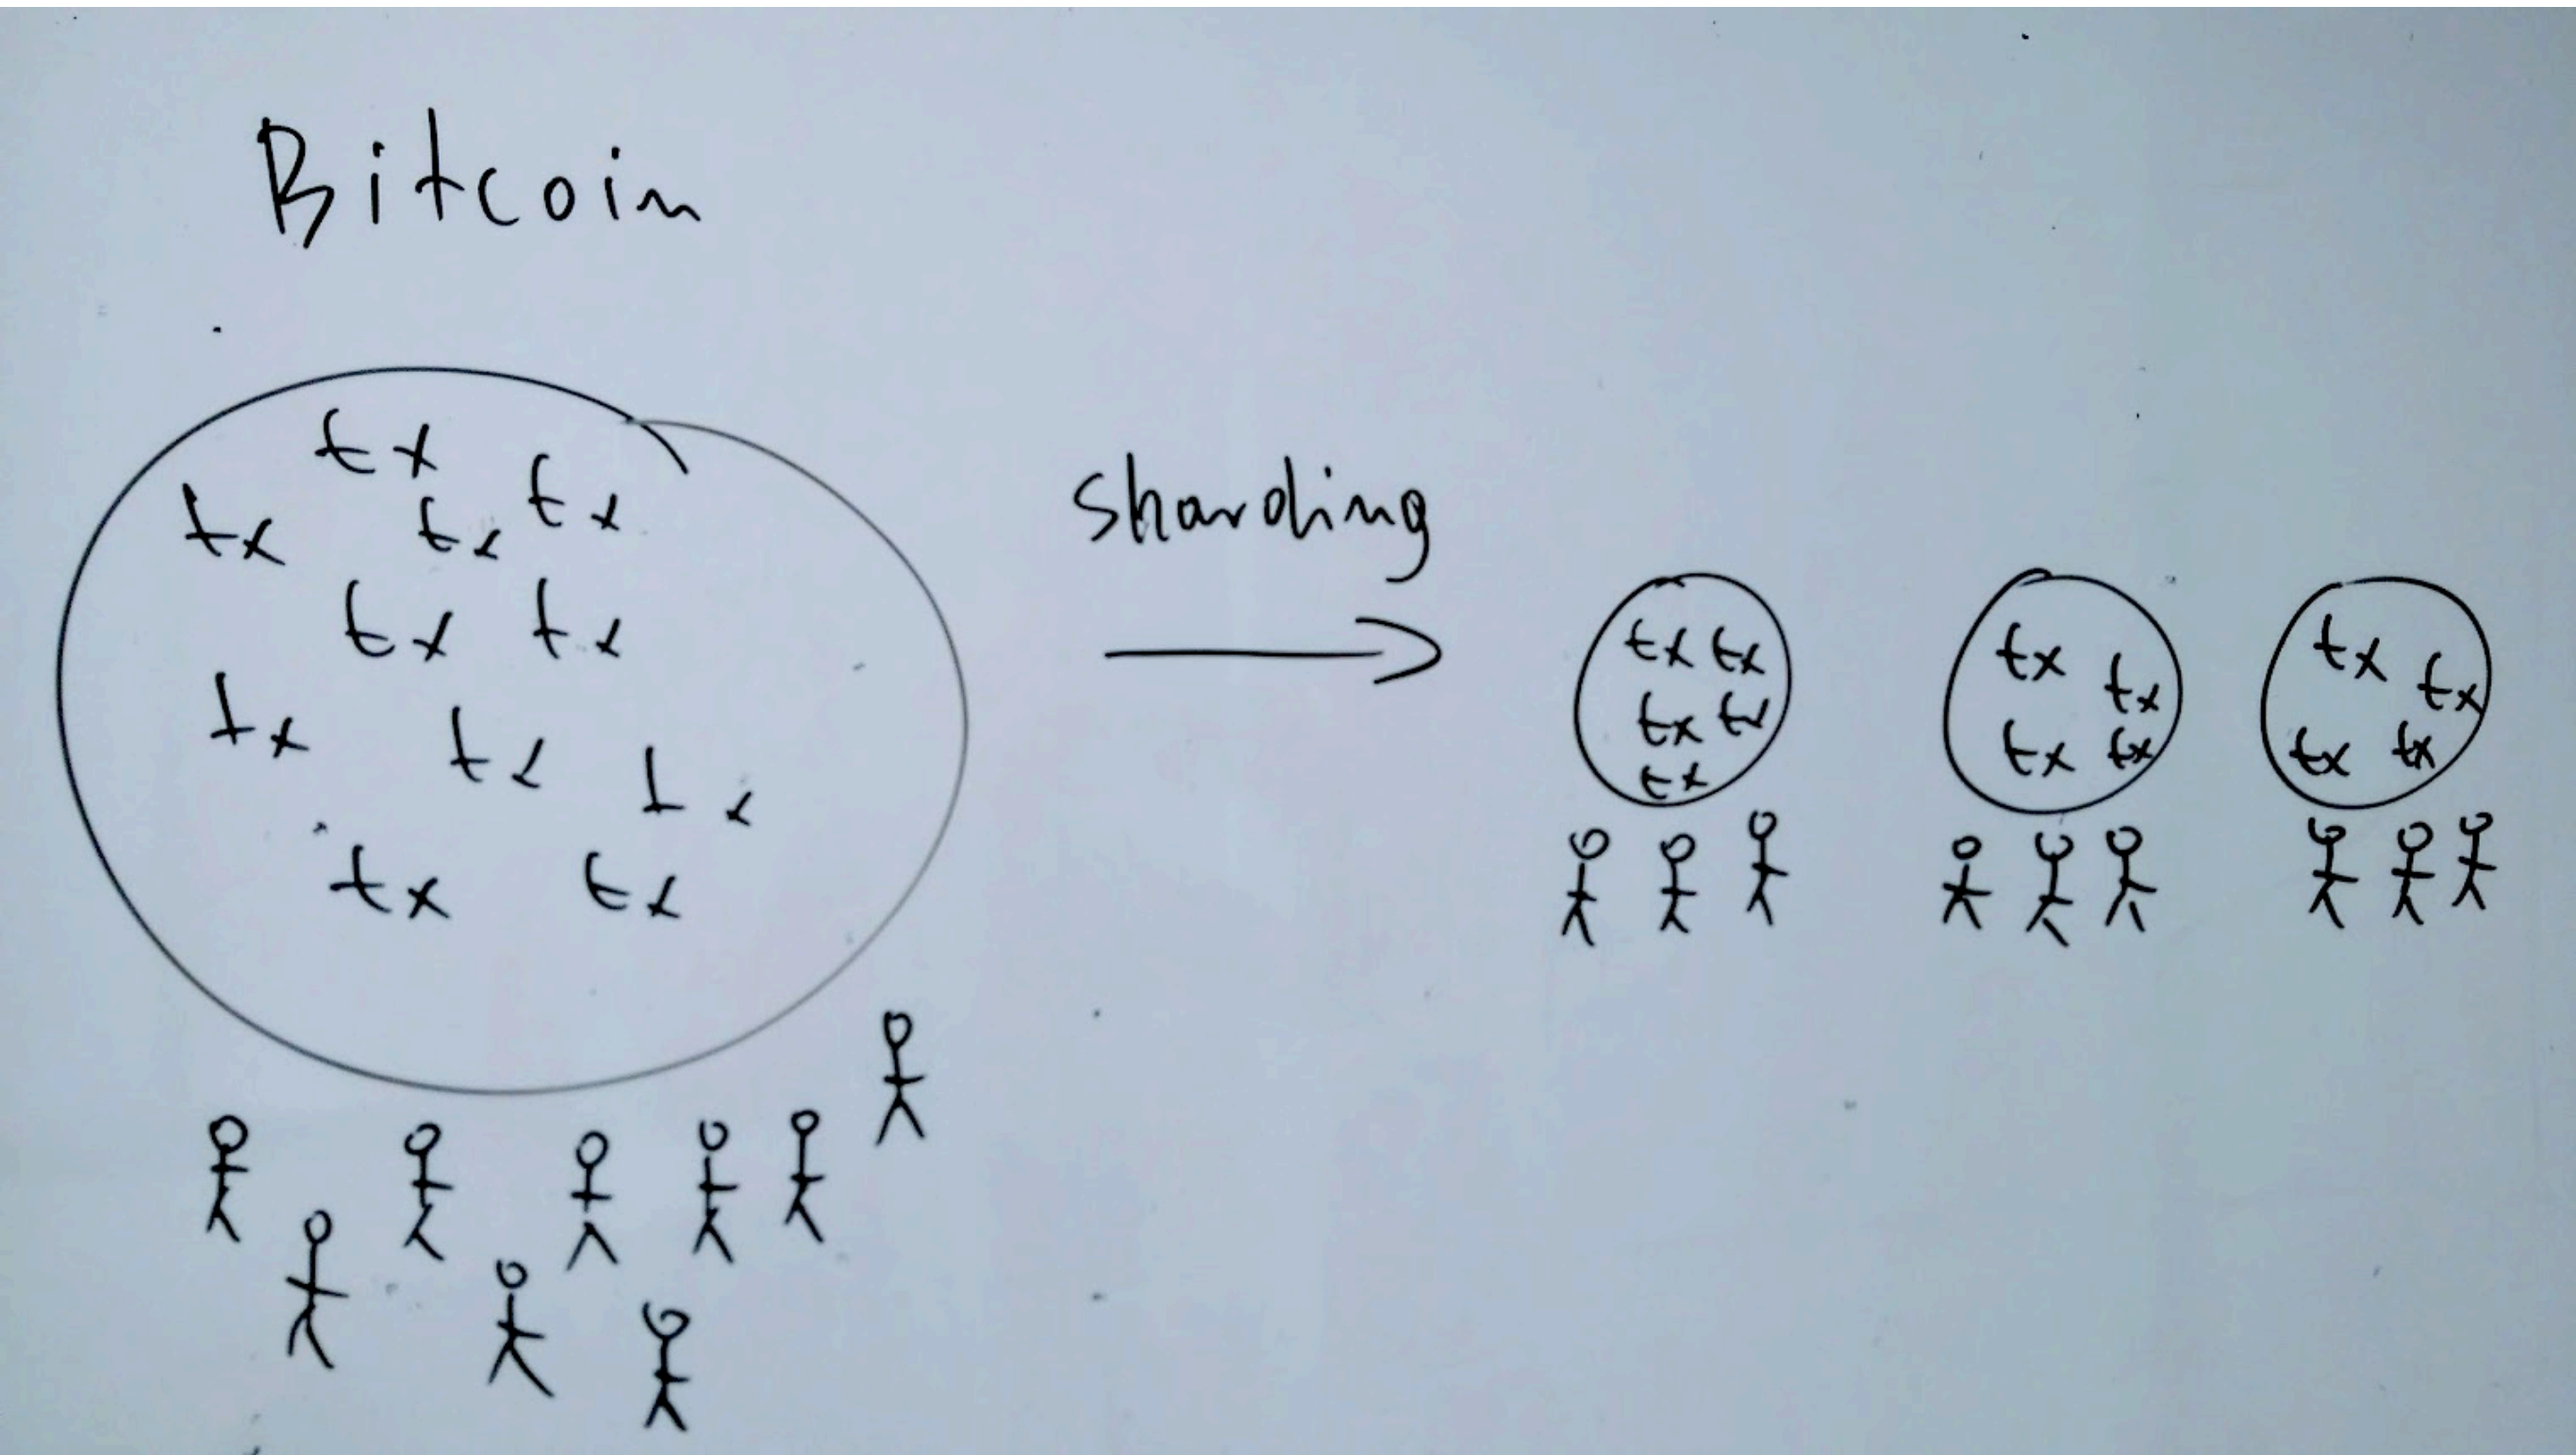

Sharding: Elastico

https://dl.acm.org/citation.cfm?id=2978389

1. Use PoW to establish identities
    - ID = H(R, IP, PubKey, Nonce) < T
    - R is security-critical, see below
2. Assign committees (use randomness of IDs)
    - Each committee has C members and a directory server (w/ members)
3. Propose a block within a committee
    - Run BFT agreement, valid blocks have 2C/3 + 1 signatures
4. Final committee to union all data blocks
    - Run BFT to produce a final block, that is then broadcast to everyone
    - R generated using R_x of final committee members

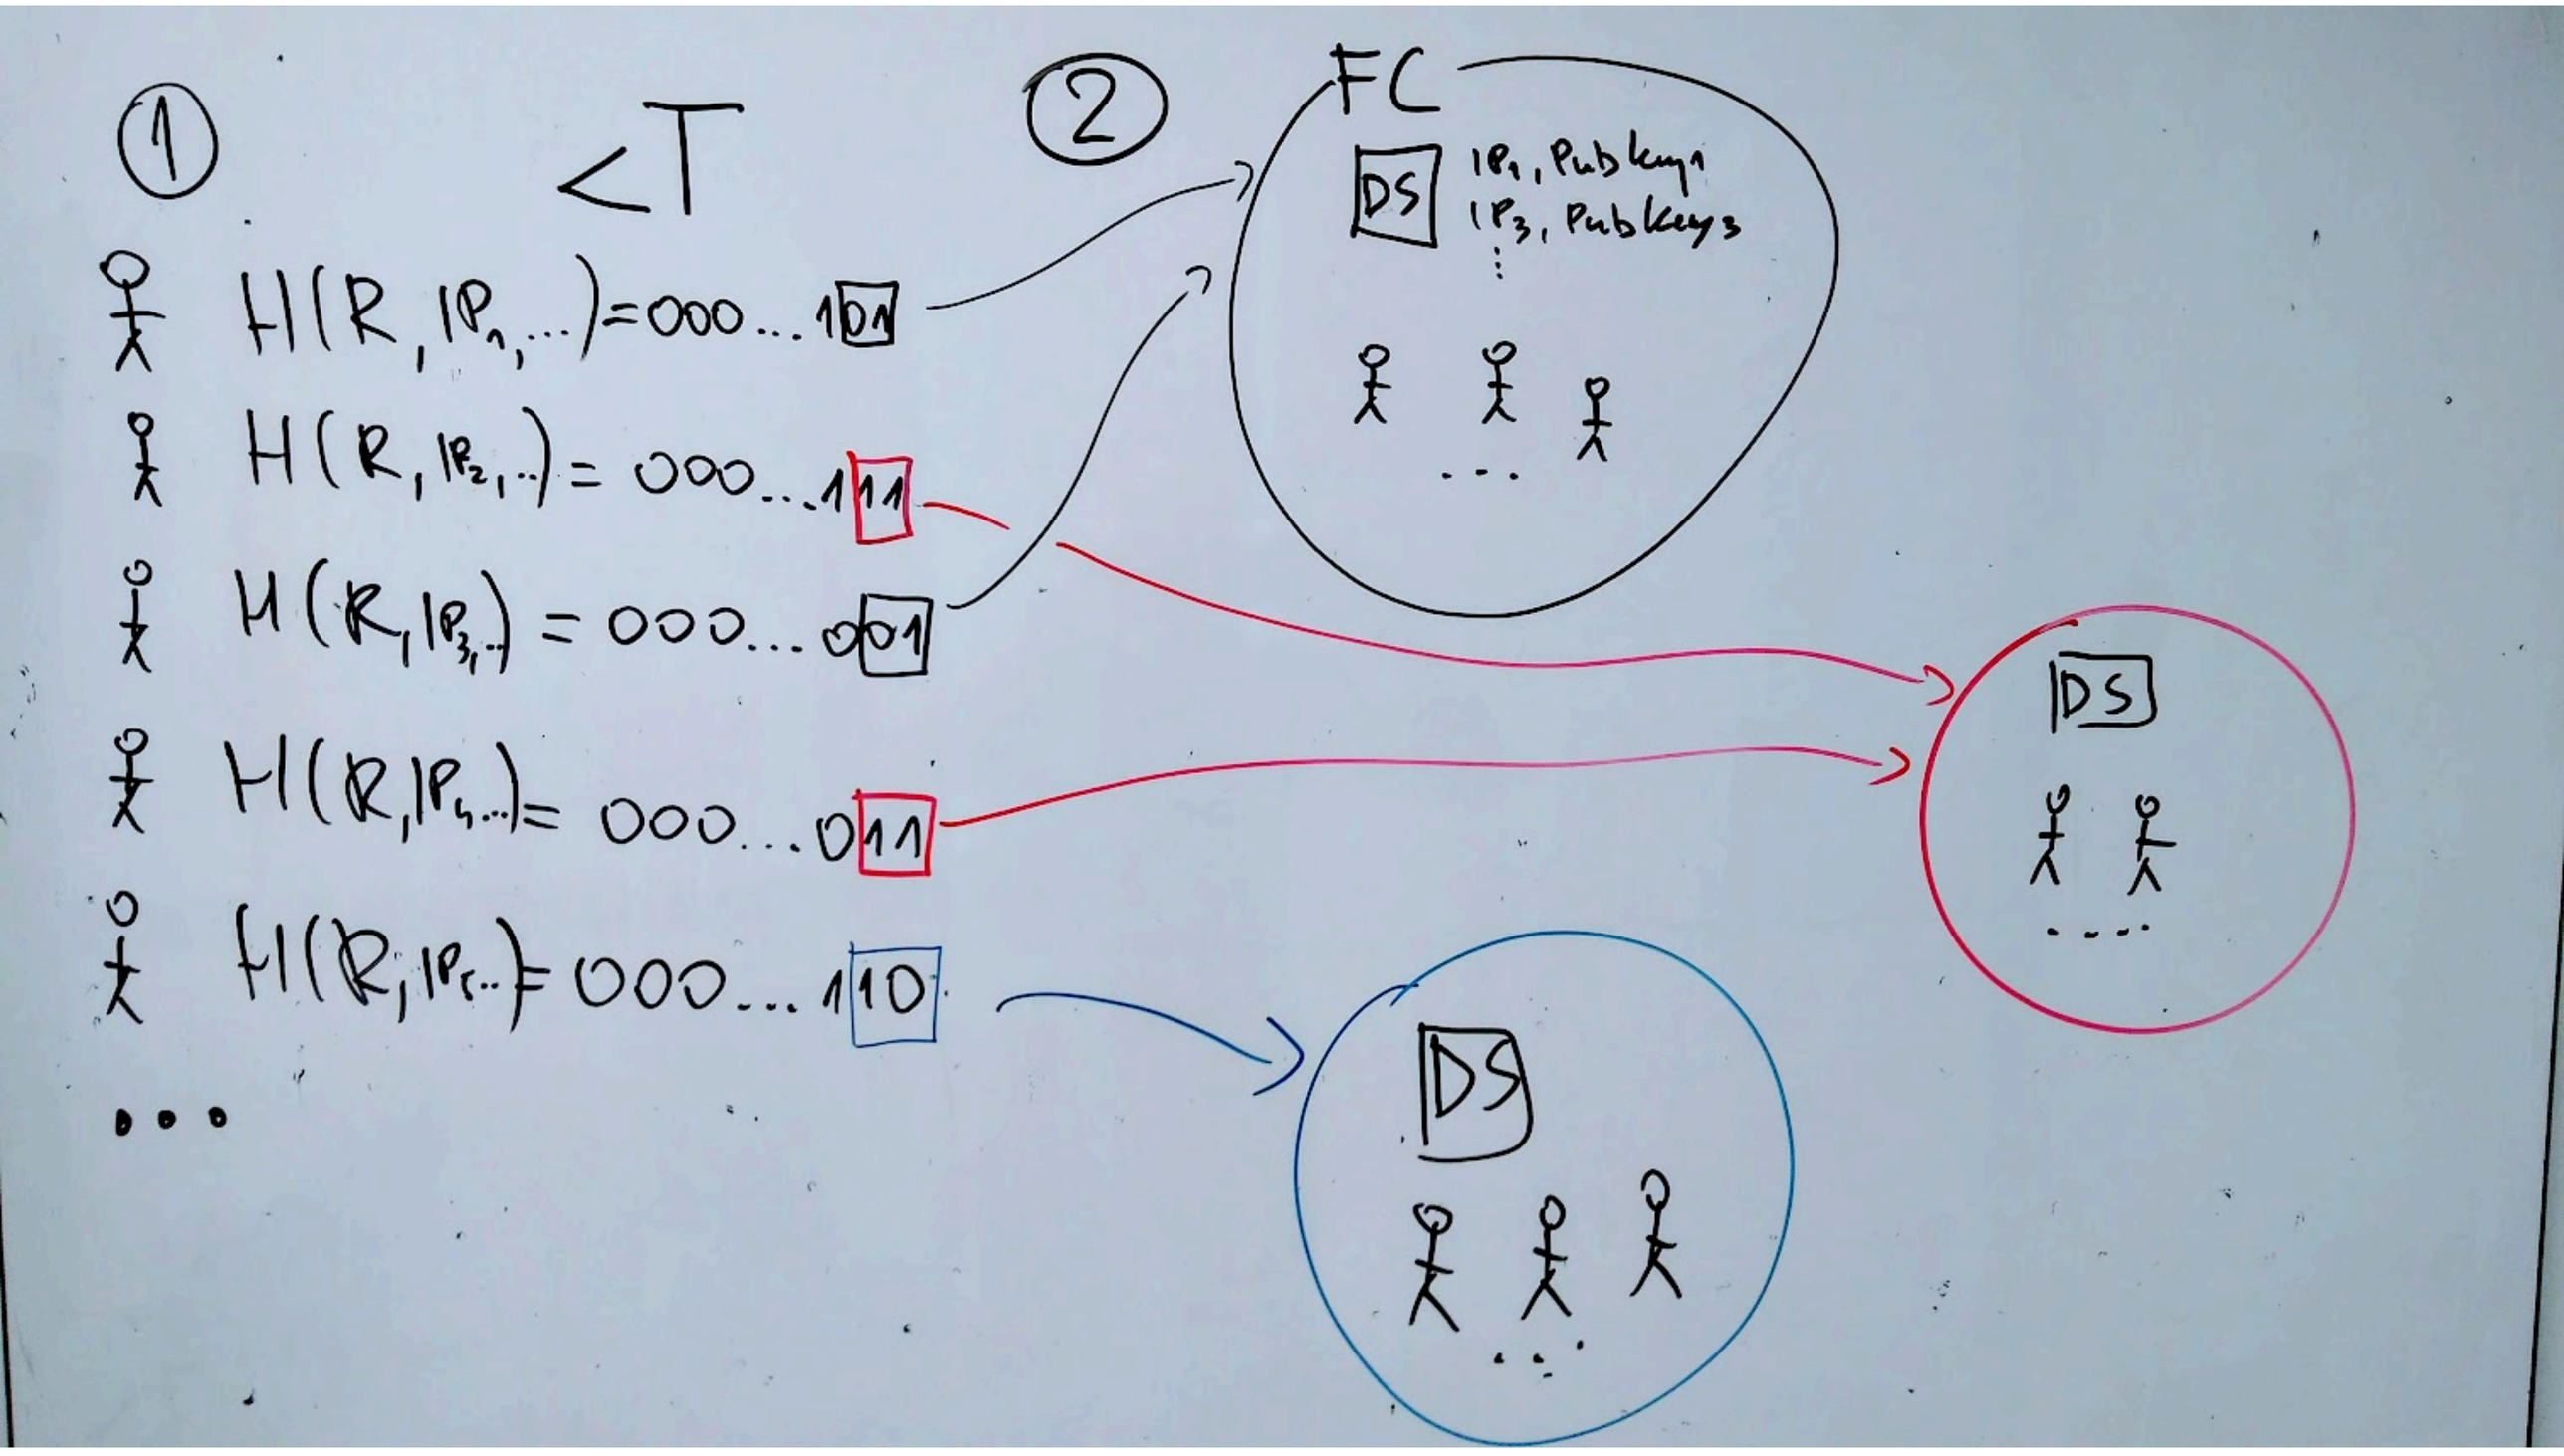
①
< T
H(R, IP1, ...) = 000...101
H(R, IP2, ...) = 000...111
H(R, IP3, ...) = 000...001
H(R, IP4, ...) = 000...011
H(R, IPc, ...) = 000...110
...
②
FC
DS
IP1, Publkey1
IP3, Publkey3
...
DS
...
DS
...

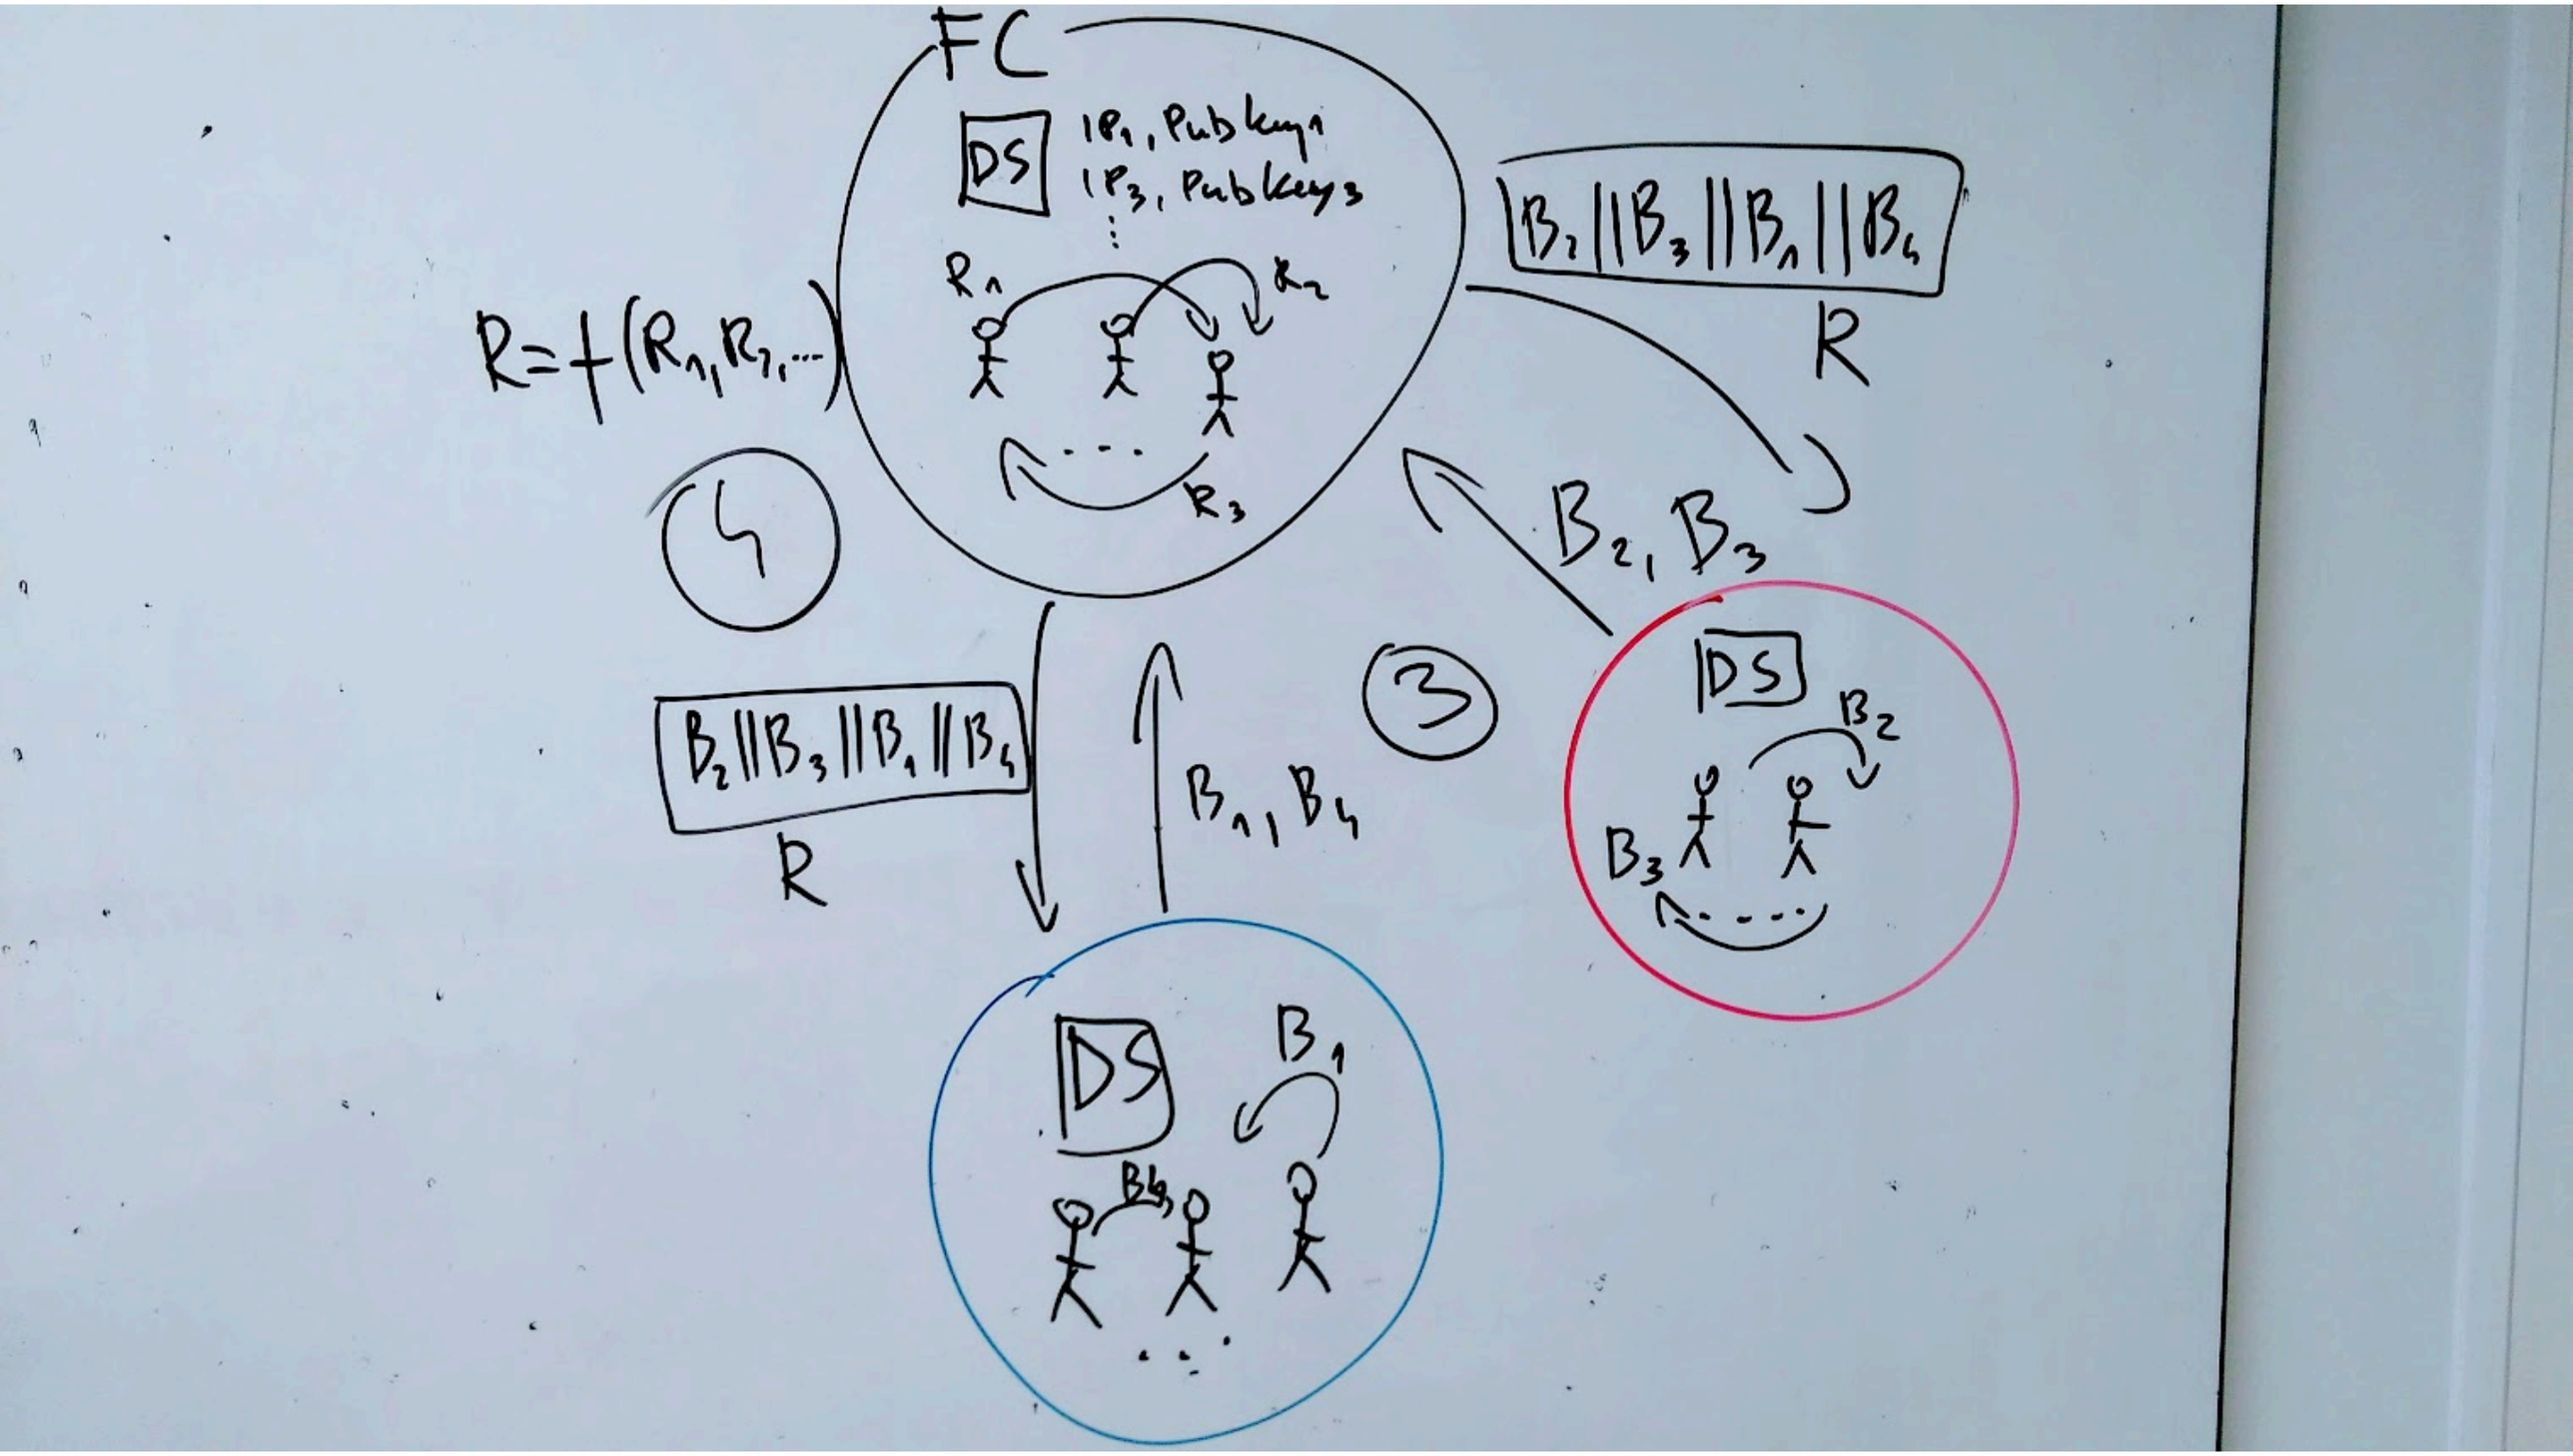
FC
DS
P_1, Pubkey1
P_3, Pubkey3
R_1
R_2
R_3
$R = f(R_1, R_2, \ldots)$
4
$B_2 || B_3 || B_1 || B_4$
R
B_2, B_3
3
B_1, B_4
$B_2 || B_3 || B_1 || B_4$
R
DS
B_2
B_3
DS
B_1
B_4

Sharding: Elastico

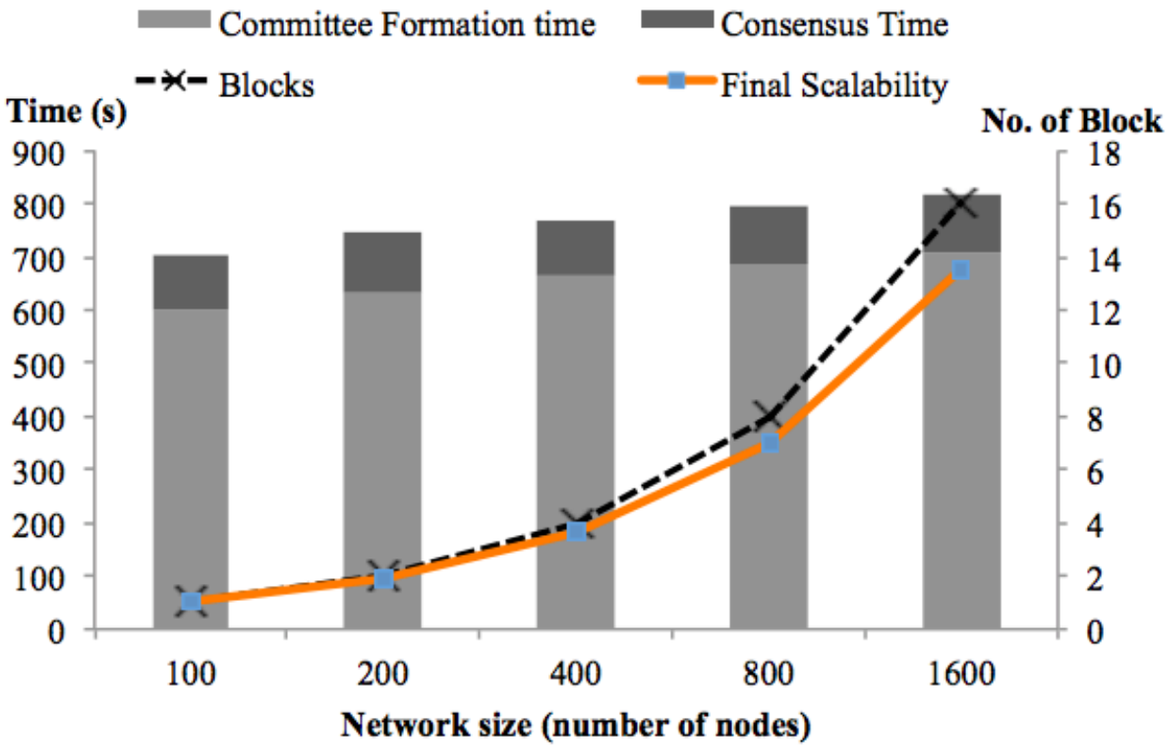

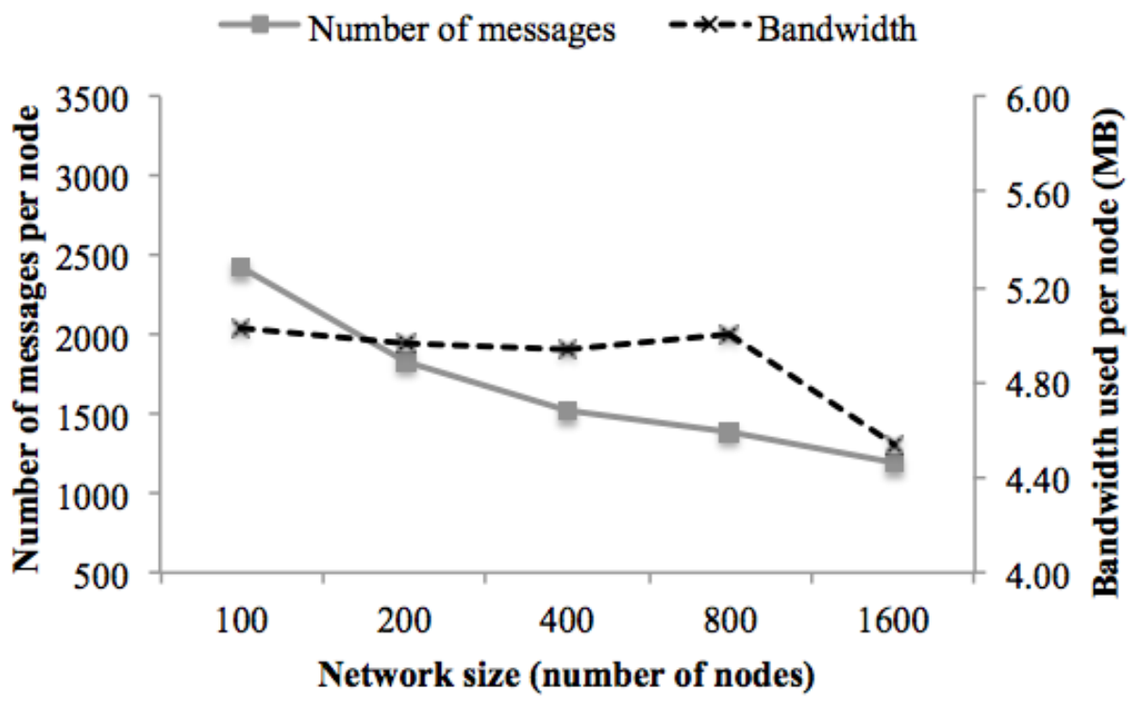

- Multiple improvements (ongoing research)
  - ZILLIQA, OmniLedger, Chainspace, Saber, …

Reading

- https://fc16.ifca.ai/bitcoin/papers/CDE+16.pdf
- https://eprint.iacr.org/2016/555.pdf
- + inline references